



Algorithms and Data Structures

Amortized Analysis

Ulf Leser

-
- Two Examples
 - Two Analysis Methods
 - Dynamic Tables
 - SOL - Analysis
-
- This lecture is not covered in [OW93] but in [Cor03]

Setting

- Note: We study a special setting - not rare, yet novel here
- We have a **sequence Q of operations** on a data structure
 - Searching and rearranging a SOL
- Operations are not independent – by changing the data structure, cost of subsequent operations is influenced
- A conventional WC-analysis produces misleading results
- **Amortized analysis** analyzes the complexity of any sequence of operations of length n
 - Or the **worst-case average cost of an operation** in any sequence
 - This is not the classical average case – we do not study an “average” sequence, but the average in any sequence

Example 1: Multi-Pop

- Assume a stack S with a special op: $\text{mpop}(k)$
- $\text{mpop}(k)$ pops $\min(k, |S|)$ elements from S
- Assume **any sequence Q** of operations
 - E.g. $Q = \{\text{push}, \text{push}, \text{mpop}(k), \text{push}, \text{push}, \text{push}, \text{mpop}(k), \dots\}$
- Assume costs $c(\text{push}) = 1$, $c(\text{pop}) = 1$, $c(\text{mpop}(k)) = k$
 - mpop simply calls pop k times
- With $|Q| = n$: What cost do we expect for Q ?
 - **Every op in Q** costs 1 (push) or 1 (pop) or k (mpop)
 - In the worst case, k can be n (n times push, than one $\text{mpop}(n)$)
 - Worst case of a single operation is $O(n)$
 - **Total worst-case cost: $O(n^2)$**



Note: Costs only $\sim 2 \cdot n$

Problem

- Clearly, the cost of Q is in $O(n^2)$, but this is **not tight**
- A simple thought shows: The cost of Q is in $O(n)$
 - Every element can be **popped only once** (no matter if this happens through a pop or a mpop)
 - Pushing an element costs 1, popping it costs 1
 - Within Q , we can **at most push $O(n)$ elements** and, hence, also only pop $O(n)$ elements
 - Thus, the **total cost is in $O(n)$**
- We want to derive such a result in a more systematic manner (analyzing SOLs is not that easy)

Example 2: Bit-Counter

- We want to generate all **bitstrings** produced by iteratively adding 1 n-times, starting from 0
- Q is a sequence of „+1“
- We count as cost of an operation the number of **bits we have to flip**
- Classical WC analysis
 - Assume bitstrings of length k
 - Roll-over counter if we exceed $2^k - 1$
 - A single operation can flip up to k bits
 - “1111111” +1
 - Worst case cost for Q: $O(k \cdot n)$

00000000		
00000001	1	1
00000010	2	3
00000011	1	4
00000100	3	7
00000101	1	8
00000110	2	10
00000111	1	11
00001000	4	15
00001001	1	16
00001010	2	18
...		

Problem

- Again, this complexity is overly pessimistic
- Cost actually is in $O(n)$
 - The right-most bit is flipped in every operation: cost= n
 - The second-rightmost bit is flipped every second time: $n/2$
 - The third ...: $n/4$
 - ...
 - Together

$$\sum_{i=0}^{k-1} \frac{n}{2^i} < n * \sum_{i=0}^{\infty} \frac{1}{2^i} = 2 * n$$

-
- Two Examples
 - Two Analysis Methods
 - Accounting Method
 - Potential Method
 - Dynamic Tables
 - SOL - Analysis

Accounting Analysis

- Idea: We **create an account** for Q
- Every operation deposits or withdraws from the account
- We choose the amounts deposited/withdrawn such that the current state of the account is always (throughout Q) an **upper bound of the actual cost of Q**
 - Let c_i be the true cost of operation i , d_i its effect on the account
 - We require

$$\forall 1 \leq k \leq n : \sum_{i=1}^k c_i \leq \sum_{i=1}^k d_i$$

- In particular, the account must never become negative
- It follows: An **upper bound for the account** is also an upper bound for the true cost

Application to mpop

- Assume $d_{\text{push}}=2$, $d_{\text{pop}}=0$, $d_{\text{mpop}}=0$
- Clearly, the account of any Q can never be zero
- We show that the sum of these costs **yields an upper bound** on the actual cost
 - Clearly, d_{push} is an upper bound on c_{push} (which is 1)
 - Whenever we push an element, we pay 1 for the actual cost and 1 for the operation that will (at same later time) **pop exactly this element**
 - It doesn't matter whether this will be in through a pop or a mpop
 - Thus, when it comes to a pop or mpop, there is **always enough money** on the account (deposited by previous push's)

- This proves:
$$\sum_{i=1}^n c_i \leq \sum_{i=1}^n d_i \leq 2 * n \in O(n)$$

Application to Bit-Counter

- Look at the **sequence Q' of flips** generated by a sequence Q
 - Every $+1$ corresponds to a sequence of 0-k flip-to-0 and **0-1 flip-to-1**
 - There is no „flip to 1“ if we roll-over
- Assume $d_{\text{flip-to-1}}=2$ and $d_{\text{flip-to-0}}=0$
 - Clearly, $d_{\text{flip-to-1}}$ is an upper bound to $c_{\text{flip-to-1}}$
 - When we flip-to-1, we pay 1 for flipping and 1 for the **back-flip-to-0 that might happen** at some later time in Q'
 - As we start with only 0 and can backflip any 1 only once, there is always enough money on the account for the flip-to-0's
 - Thus, the account is an upper bound on the actual cost
- As **every operation in Q can pay at most 2** (there is at most 1 flip-to-1), Q is in $O(n)$

-
- Two Examples
 - Two Analysis Methods
 - Accounting Method
 - Potential Method
 - Dynamic Tables
 - SOL - Analysis

Potential Method: Idea

- In the accounting method, we assign a cost to every operation and compare aggregated accounting costs of ops with real costs of ops
- In the potential method, we assign a **potential $\Phi(D)$ to the data structure D** manipulated by Q
- As ops from Q change D , they also change D 's potential
- The trick is to design Φ such that we can (again) use it to derive an **upper bound on the real cost** of Q

Potential Function

- Let D_0, D_1, \dots, D_n be the states of D when applying Q
- We define the **amortized cost** d_i of the i 'th operation as $d_i = c_i + \Phi(D_i) - \Phi(D_{i-1})$
- We derive the **amortized cost of Q** as

$$\sum_{i=1}^n d_i = \sum_{i=1}^n (c_i + \phi(D_i) - \phi(D_{i-1})) = \sum_{i=1}^n c_i + \phi(D_n) - \phi(D_0)$$

- If we find a Φ such that (a) we get formulas for the amortized costs of all d_i and (b) $\Phi(D_n) \geq \Phi(D_0)$, we have an **upper bound for the real costs**

Always Pay in Advance

- Operations raise or lower the potential of D
- We need to find a function Φ such that
 - 1: $\Phi(D_i)$ depends on a property of D
 - 2: $\Phi(D_n) \geq \Phi(D_0)$ [and we will always have $\Phi(D_0) = 0$]
 - 3: We can compute $d_i = c_i + \Phi(D_i) - \Phi(D_{i-1})$ for any possible op
- As within a sequence we do not know its future, we also have to require that $\Phi(D_i)$ never is negative
 - Otherwise, the amortized cost of the sequence $Q[1..i]$ is no upper bound in the real costs
- We need to pay in advance

Example: mpop

- We use the **number of objects on the stack** as its potential
- Then
 - 1: $\Phi(D_i)$ depends on a property of D
 - 2: $\Phi(D_n) \geq \Phi(D_0)$ and $\Phi(D_0) = 0$
 - 3: Compute $d_i = c_i + \Phi(D_i) - \Phi(D_{i-1})$
 - If op is push: $d_i = c_i + 1 = 2$
 - If op is pop: $d_i = c_i - 1 = 0$
 - If op is mpop(k): $d_i = c_i - k = 0$
- Thus, **$2*n \geq \sum d_i \geq \sum c_i$** and Q is in $O(n)$

Example: Bit-Counter

- We use the number of „1“ in the bitstring as its potential
- Then
 - 1: $\Phi(D_i)$ depends on a property of D
 - 2: $\Phi(D_n) \geq \Phi(D_0)$ and $\Phi(D_0) = 0$
 - 3: Compute $d_i = c_i + \Phi(D_i) - \Phi(D_{i-1})$
 - Let the i 'th operation incur t_i flip-to-0 and 0 or 1 flip-to-1
 - Thus, $c_i \leq t_i + 1$
 - If $\Phi(D_i) = 0$, then this op has flipped all positions to 0, and previously they were all 1 and we have $\Phi(D_{i-1}) = k$
 - If $\Phi(D_i) > 0$, then $\Phi(D_i) = \Phi(D_{i-1}) - t_i + 1$
 - In both cases, we have $\Phi(D_i) \leq \Phi(D_{i-1}) - t_i + 1$
 - Thus, $d_i = c_i + \Phi(D_i) - \Phi(D_{i-1}) \leq (t_i + 1) + (\Phi(D_{i-1}) - t_i + 1) - \Phi(D_{i-1}) \leq 2$
- Thus, $2 * n \geq \sum d_i \geq \sum c_i$ and Q is in $O(n)$

-
- Two Examples
 - Two Analysis Methods
 - **Dynamic Tables**
 - SOL - Analysis

Dynamic Tables

- We now use amortized analysis for something more useful: Complexity of operations on a **dynamic table**
- Assume an array T and a sequence **Q of insert/delete ops**
- We want to keep the array small, yet avoid overflows
- Method: When inserting and T is full, **we double $|T|$** ; upon deleting and A is only half-full, we reduce $|T|$ by 50%
 - Actually, deletion will be handled differently later
- Thus, if the i 'th operation is a insertion (or deletion), it costs **either 1 or i** ; as i can be up to n , the complexity of insertion is $O(n)$; thus, the complexity of Q is $O(n^2)$
 - Assuming that we copy A to a new location for growing / shrinking

Amortized Analysis with Potential Method

- Let $\text{num}(T)$ be the current number of elements in T
- We use potential $\Phi(T) = 2 * \text{num}(T) - |T|$
 - 1: Of course
 - 2: As T is always at least half-full, $\Phi(T)$ is always ≥ 0
 - 2: We start with $|T|=0$, and thus $\Phi(T_n) - \Phi(T_0) \geq 0$
 - Intuitively a potential
 - Immediately before an expansion, $\text{num}(T)=|T|$ and $\Phi(T)=|T|$, so there is **much potential in T** (we saved for the expansion to come)
 - Immediately after an expansion, $\text{num}(T)=|T|/2$ and $\Phi(T)=0$; **all potential has been used**, we need to save again for the next expansion

Operations

- 3: Let's study $d_i = c_i + \Phi(T_i) - \Phi(T_{i-1})$ for insertions

- Without expansion

$$\begin{aligned}d_i &= 1 + (2 * \text{num}(T_i) - |T_i|) - (2 * \text{num}(T_{i-1}) - |T_{i-1}|) \\&= 1 + 2 * \text{num}(T_i) - 2 * \text{num}(T_{i-1}) - |T_i| + |T_{i-1}| \\&= 1 + 2 + 0 \\&= 3\end{aligned}$$

- With expansion

$$\begin{aligned}d_i &= \text{num}(T_i) + (2 * \text{num}(T_i) - |T_i|) - (2 * \text{num}(T_{i-1}) - |T_{i-1}|) \\&= \text{num}(T_i) + 2 * \text{num}(T_i) - 2 * (\text{num}(T_i) - 1) - 2 * (\text{num}(T_i) - 1) \\&\quad \quad \quad + \text{num}(T_i) - 1 \\&= 3 * \text{num}(T_i) - 2 * \text{num}(T_i) + 2 - 2 * \text{num}(T_i) + 2 + \text{num}(T_i) - 1 \\&= 3\end{aligned}$$

- Thus, $3 * n \geq \sum d_i \geq \sum c_i$ and Q is in $O(n)$ (for only insertions)

Why 3

- There is an intuitive explanation for the **cost of exactly 3**
- Think in terms of accounting method
- When we insert an element, we deposit 3 on the account
 - 1 for the insertion (the real cost)
 - 1 for the time when we need to **copy this new element** at the next expansion
 - These 1's fill the account with $|T_i|/2$ before the next expansion
 - 1 for **one of the $|T_i|/2$ elements** already in A
 - These fill the account with $|T_i|/2$ before the next expansion
- Thus, we have enough credit at the next expansion

Problem: Deletions

- Our strategy for deletions so far is not very clever
 - Assume a table with $\text{num}(T) = |T|$
 - Assume a sequence $Q = \{I, D, I, D, I, D, I \dots\}$
 - This sequence will perform $|T| + |T|/2 + |T| + |T|/2 + \dots$ real ops
 - As $|T|$ is $O(n)$, Q is in $O(n^2)$ and not in $O(n)$
- Simple trick: Do only contract when $\text{num}(T) = |T|/4$
 - Leads to amortized cost of $O(n)$ for any sequence of operations
 - We omit the proof (see [Cor03])

-
- Two Examples
 - Two Analysis Methods
 - Dynamic Tables
 - SOL – Analysis
 - Goal and idea
 - Preliminaries
 - A short proof

Re-Organization Strategies

- Think of self-organizing lists again
- When searching an element, we **change the list L**
 - As usual: Accessing the i 'th element costs i
- Three popular strategies
 - **MF, move-to-front:**
After searching an element e , move e to the front of L
 - **T, transpose:**
After searching an element e , swap e with its predecessor in L
 - **FC, frequency count:**
Keep an access frequency counter for every element in L and keep L sorted by this counter. After searching e , increase counter of e and move “up” to keep sorted'ness

Notation

- Assume we have a **self-organizing strategy A** and a sequence $S = \{s_i\}$ of accesses to L
- After an access to element i , **A may move i by swapping**
 - Swap with predecessor (to-front) or successor (to-back)
 - Let $F_A(l)$ be the number of front-swaps and $X_A(l)$ the number of back-swaps after access number l
 - F_A/X_A for strategy A, F_{MF}/X_{MF} for strategy MF, $F_T/X_T \dots F_{FC}/X_{FC}$
 - Of course, $\forall l: X_{MF}(l) = X_T(l) = X_{FC}(l) = 0$
- Let **$C_A(S)$ be the total access costs** of A incurred by S
 - Again: C_{MF} for strategy MF, C_T for T, C_{FC} for FC
- With conventional worst-case analysis, we can only derive that $C_A(S)$ is in $O(|S| * |L|)$ – for any strategy

Theorem

- Theorem (Amortized costs)

*Let A be **any self-organizing strategy** for a SOL L , MF be the move-to-front strategy, and S be a sequence of accesses to L . Then*

$$C_{MF}(S) \leq 2 * C_A(S) + X_A(S) - F_A(S) - |S|$$

- What does this mean?
 - We don't learn more about the absolute complexity of A / MF
 - But we learn that **MF is quite good**
 - Any strategy following the same constraints (only series of swaps) will at best be **roughly twice as good as MF**
 - Assuming $C(S) \gg |S|$ and for $|S| \rightarrow \infty$: $X(S) \sim F(S)$ for any strategy
 - Despite its simplicity, MF is a fairly safe bet in whatever circumstances (= sequences)

Idea of the Proof

- We will compare access costs in L using MF and A
- Think of both strategies **running S on two copies** of the same initial list L
- After each step, A and MF perform different swaps, so all list states except the first very likely are different
- We will compare list states and determine a certain property – the **number of inversions** (“Fehlstellungen”)
 - Actually, we only analyze how the number of inversion changes
- We will show that the **number of inversions define a potential** of a pair of lists and can be used to derive **an upper bound on the differences in real costs**

Content of this Lecture



- Two Examples
- Two Analysis Methods
- Dynamic Tables
- SOL - Analysis
 - Goal and idea
 - Preliminaries
 - A short proof

Inversions

- Let L and L' be permutation of the set $\{1, 2, \dots, n\}$
- Definition
 - A pair (i, j) is called an *inversion of L and L'* iff i and j are in different order in L than in L' (for $0 \leq i, j \leq n$ and $i \neq j$)
 - The number of inversions between L and L' is written $\text{inv}(L, L')$
- Remarks
 - Different order: Once i before j , once i after j
 - Obviously, $\text{inv}(L, L') = \text{inv}(L', L)$
- Example: $\text{inv}(\{4, 3, 1, 5, 7, 2, 6\}, \{3, 6, 2, 5, 1, 4, 7\}) = 12$
- Without loss of generality, we assume that $L = \{1, \dots, n\}$
 - Because we only look at changes in number of inversions and not at the actual set of inversions



Sequences of Changes

- Assume we applied $l-1$ steps creating L_{MF} using MF and L_A using A
- Let us consider the **next step l** , creating L_{MF}' and L_A'




Inversion Changes

- How does **I** change the number of inversions between L_{MF} and L_A ; can we compute $\text{inv}(L_{MF}', L_A')$ from $\text{inv}(L_{MF}, L_A)$?
 - Assume **I** accesses element i from L_A
 - We may assume it is at position i
 - Let i be at position k in L_{MF}
 - Access in L_A costs i , in L_{MF} it costs k
 - After **I**, **A** performs an unknown number of swaps; **MF** performs exactly $k-1$ front-swaps




Counting Inversion Changes 1

- Let X_i be the set of values that are **before position k in L_{MF} and after position i in L_A**
- Let Y_i be the values before position k in L_{MF} and before i in L_A
 - Clearly, $|X_i| + |Y_i| = k-1$
- All pairs (i, c) with $c \in X_i$ are inversions between L_A and L_{MF}
 - There may be more; but only those including i are affected
- After step 1, **MF moves element i to the front**
 - All inversions (i, c) with $c \in X_i$ disappear (there are $|X_i|$ many)
 - But $|Y_i| = k-1 - |X_i|$ new inversions appear
- If A did nothing: $\text{inv}(L_{MF}', L_A') = \text{inv}(L_{MF}, L_A) - |X_i| + k-1 - |X_i|$
 - But **A is doing something**



Counting Inversion Changes 2

- In step I, let A perform $F_A(i)$
front-swaps and $X_A(i)$
back-swaps



- Every front-swap (swapping i before j) in L_A decreases $\text{inv}(L_{MF}', L_A')$ by 1
 - Before the swap, j must be before i in L_A (it is a front-swap), but after i in L_{MF}' (because i now is the first element in L_{MF}')
 - After the swap, i is before j in both L_A' and L_{MF}'
- Equally, every back-swap increases $\text{inv}(L_{MF}', L_A')$ by 1
- Together: After step I, we have

$$\text{inv}(L_{MF}', L_A') = \underbrace{\text{inv}(L_{MF}, L_A)}_{\text{Before step I}} - \underbrace{|X_i|}_{\text{New by MF}} + \underbrace{k-1-|X_i|}_{\text{New by A}} - \underbrace{F_A(i)}_{\text{New by A}} + \underbrace{X_A(i)}_{\text{New by A}}$$

Amortized Costs

- Let t_i be the real costs of strategy MF for step i
- We use the number of inversions $\Phi(L_A, L_{MF}) = \text{inv}(L_A^i, L_{MF}^i)$ as the **potential function of the pair of data structures L_A, L_{MF}**
- Definition
 - The **amortized costs of step i , called a_i** , are

$$a_i = t_i + \text{inv}(L_A^i, L_{MF}^i) - \text{inv}(L_A^{i-1}, L_{MF}^{i-1})$$
 - Accordingly, the amortized costs of sequence S , $|S|=m$, are

$$\sum a_i = \sum t_i + \text{inv}(L_A^m, L_{MF}^m) - \text{inv}(L_A^0, L_{MF}^0)$$
- This is a proper potential function
 - 1: Φ depends on a property of the pair L_A, L_{MF}
 - 2: $\text{inv}()$ can never be negative, so $\Phi(L_A^n, L_{MF}^n) \geq \Phi(L, L) = 0$
- Let's look at **how operations change the potential**

Content of this Lecture

- Two Examples
- Two Analysis Methods
- Dynamic Tables
- SOL - Analysis
 - Goal and idea
 - Preliminaries
 - A short proof (after much preparatory work)

Putting it Together

- We know for every step l from S accessing i :
 $\text{inv}(L_{MF'}, L_{A'}) = \text{inv}(L_{MF}, L_A) - |X_i| + k-1-|X_i| - F_A(i) + X_A(i)$
and thus
 $\text{inv}(L_{MF'}, L_{A'}) - \text{inv}(L_{MF}, L_A) = -|X_i| + k-1-|X_i| - F_A(i) + X_A(i)$
- Using the fact that $t_l = k$ for MF , we get **amortized costs** of
$$\begin{aligned} a_l &= t_l + \text{inv}(L_{A'}, L_{MF'}) - \text{inv}(L_A, L_{MF}) \\ &= k - |X_i| + k-1-|X_i| - F_A(i) + X_A(i) \\ &= 2(k-|X_i|) - 1 - F_A(i) + X_A(i) \end{aligned}$$
- Recall that $|Y_i| = k-1-|X_i|$ are those elements before i in both lists. This implies that $k-1-|X_i| \leq i-1$ or $k-|X_i| \leq i$
 - There can be at most $i-1$ elements before position i in L_A
- Therefore: $a_l \leq 2i - 1 - F_A(i) + X_A(i)$

Putting it Together

- This is the **central trick!**
- Because we only looked at inversions (and hence the sequence of values), we can draw a connection between the value that is accessed and the number of inversions that are affected
- Recall that $|Y_i| = k - 1 - |X_i|$ are those elements before i in both lists. This implies that $k - 1 - |X_i| \leq i - 1$ or $k - |X_i| \leq i$
 - There can be at most $i - 1$ elements before position i in L_A
- Therefore: $a_i \leq 2i - 1 - F_A(i) + X_A(i)$

Aggregating

- We also know the cost of accessing i using A : that's i
- Together: $a_i \leq 2C_A(i) - 1 - F_A(i) + X_A(i)$
- Aggregating this inequality over all a_i (hence S), we get

$$\sum a_i \leq 2 * C_A(S) - |S| - F_A(S) + X_A(S)$$

- By definition, we also have

$$\sum a_i = \sum t_i + \text{inv}(L_A^m, L_{MF}^m) - \text{inv}(L_A^0, L_{MF}^0)$$

- Since $\sum t_i = C_{MF}(S)$ and $\text{inv}(L_A^0, L_{MF}^0) = 0$, we get
$$C_{MF}(S) + \text{inv}(L_A^m, L_{MF}^m) \leq 2 * C_A(S) - |S| - F_A(S) + X_A(S)$$
- It finally follows ($\text{inv}() \geq 0$)

$$C_{MF}(S) \leq 2 * C_A(S) - |S| - F_A(S) + X_A(S)$$

Summary

- Self-organization creates a type of problem we were not confronted with before
 - Things change during program execution
 - But not at random – we follow a strategy
- Analysis is none-trivial, but
 - Helped to find a elegant and surprising conjecture
 - Very interesting in itself: We showed relationships between measures we never counted (and could not count easily)
 - Original work: Sleator, D. D. and Tarjan, R. E. (1985). "Amortized efficiency of list update and paging rules." *Communications of the ACM* 28(2): 202-208.