



Algorithms and Data Structures


Priority Queues

Ulf Leser

Special Scenarios for Searching


- Up to now, we assumed that all elements of a list are equally important and that any of them could be searched next (with varying probability)
- What if some elements are **more important** than others?
 - There is a (maybe partial) order on list elements
 - The most important elements are always retrieved next
 - **Priority Queues**
- Difference to Self-Organizing Lists
 - Most important element is **always retrieved next** – should be $O(1)$
 - List must be kept ordered by importance
 - We look at a scenario where new **elements are inserted** all the time and the **most important element is removed** regularly

Shortest Paths in a Graph



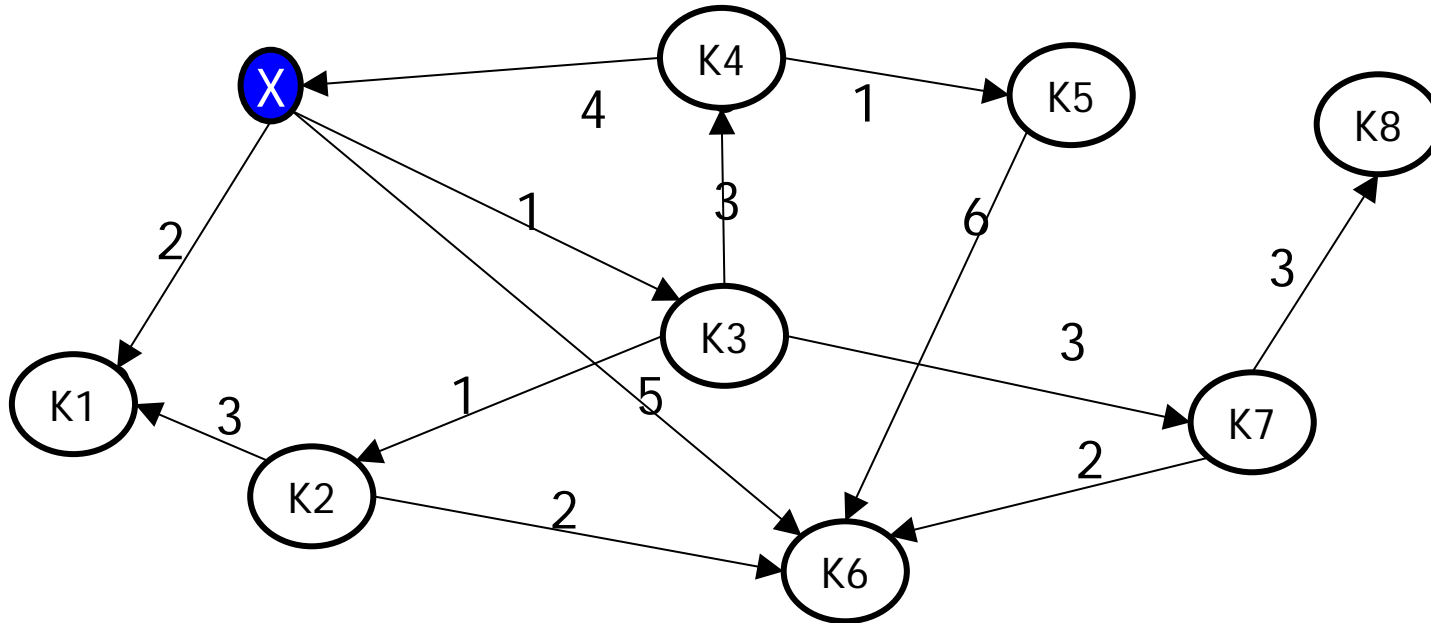
- Task: Find the **distance between X** and all other nodes
 - Classical problem: Single-Sink-Shortest-Paths
 - Famous solution: **Dijkstra's algorithm**

Assumptions




- We assume that there is at least one path between X and any other node (every node is reachable from X)
- We assume strictly positive edge weights
- Distance is the length (=sum of weights) of the shortest path
- There might be many shortest paths, but distance is unique
- We only want the distance and need no “witness path”

Exhaustive Solution




- First approach: **Enumerate all paths**
 - Need to **break cycles** (e.g. $X - K3 - K4 - X - K3 - \dots$)
 - Using DFS: $X - K3 - K4 - X$ [BT-K4] $- K5 - K6$ [BT-K5] [BT-K4] [BT-K3] $- K7 - K8$ [BT-K7] $- K6$ [BT-K7] [BT-K3] $- K2 - K6$ [BT-K2] $- K1$ [BT-K2] [BT-K3] [BT-X] $- K6 - \dots$

Redundant work



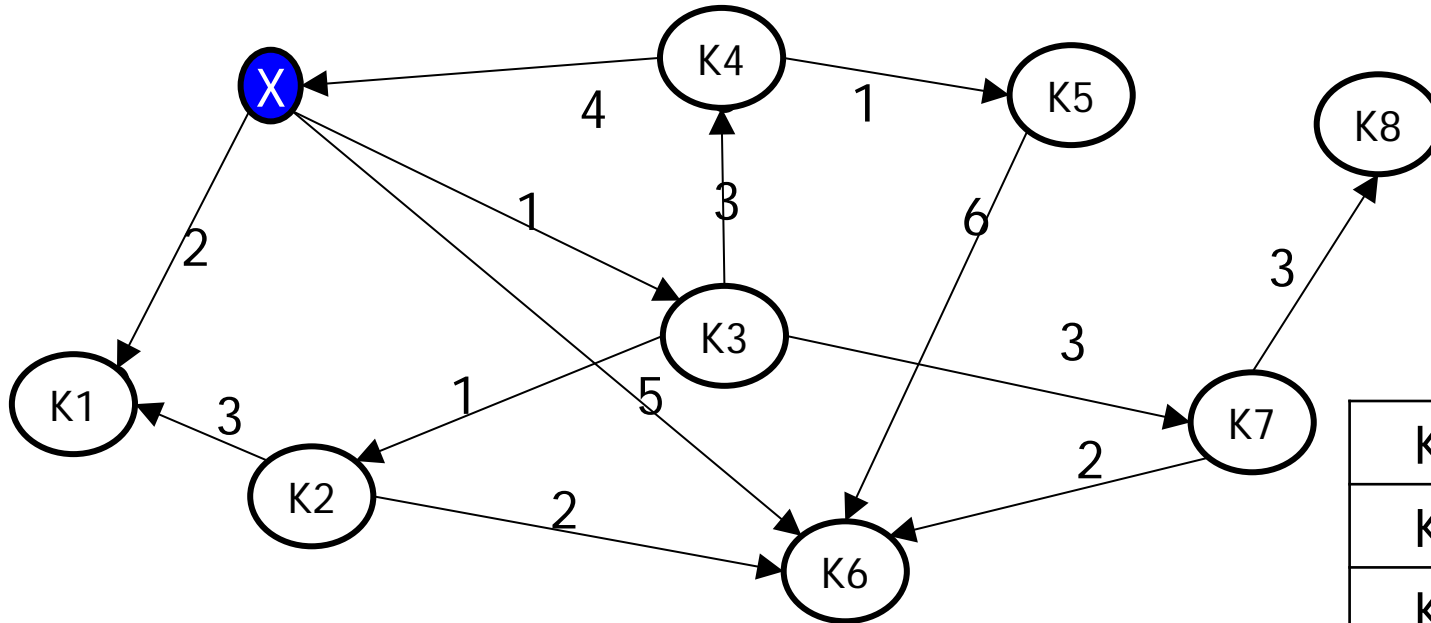
- First approach: Enumerate all paths
 - Need to break cycles (e.g. $X - K3 - K4 - X - K3 - \dots$)
 - Using DFS: $X - K3 - K4 - X$ [BT-K4] $- K5 - K6$ [BT-K5] [BT-K4] [BT-K3] $- K7 - K8$ [BT-K7] $- K6$ [BT-K7] [BT-K3] $- K2 - K6$ [BT-K2] $- K1$ [BT-K2] [BT-K3] [BT-X] $K6 - \dots$

Dijkstra's Idea



- Enumerate **paths by their length** (neither DFS nor BFS)
- Assume we reach a node **Y** by a **path p** of **length l** and we have already explored all paths with length $l' \leq l$ and that Y was not reached yet
 - We always mean “all paths starting from X”
- Then **p must be the shortest path** between X and Y
 - Because any p' between X and Y would have a **prefix of length at least l** and (a) a continuation with length > 0 or (b) would not need a continuation (then p is as short as p')

Example for Idea



- 1: X – K3
- 2: X – K3 – K2
- 2: X – K1
- 4: X – K3 – K2 – K6
- 4: X – K3 – K4
- 4: X – K3 – K7

- 5: X – K3 – K4 – K5
- 7: X – K3 – K7 – K8
- Stop (all nodes found)

| | |
|----|---|
| K3 | 1 |
| K2 | 2 |
| K1 | 2 |
| K6 | 4 |
| K4 | 4 |
| K7 | 4 |
| K5 | 5 |
| K8 | 7 |

A Further Trick


- We enumerate paths by length by iteratively extending short paths by all **possible next edges**
 - I.e., by looking at all edges outgoing from the end node of a short path
- These extensions
 - ... either lead to a node which we didn't reach yet – then we found a path, but cannot yet be sure it is the shortest
 - ... or lead to a node which we already reached but we are not yet sure of we found the shortest path to it – **update current best distance**
 - ... or lead to a node which we already reached and for which we also surely found the a shortest path already – these can be ignored
- Eventually, we can **enumerate nodes by their distance**

Algorithm

```
1. G = (V, E);
2. x : start_node;      # x ∈ V
3. A : array_of_distances;
4. ∀i: A[i] := ∞;
5. L := V;
6. A[x] := 0;
7. while L ≠ ∅
8.   k := L.get_closest_node();
9.   L := L \ k;
10.  forall (k, f, w) ∈ E do
11.    if f ∈ L then
12.      new_dist := A[k] + w;
13.      if new_dist < A[f] then
14.        A[f] := new_dist;
15.      end if;
16.    end if;
17.  end for;
18. end while;
```

- Assumptions
 - Nodes have IDs between 1 ... |V|
 - Edges are (from, to, weight)
- We enumerate nodes by length of their shortest paths
 - In the first loop, we pick x and update distances (A) to all adjacent nodes
 - When we pick a node k, we already have computed its distance to x in A
 - We adapt the current best distances to all neighbors of k we haven't picked yet
- Once we picked all nodes, we are done

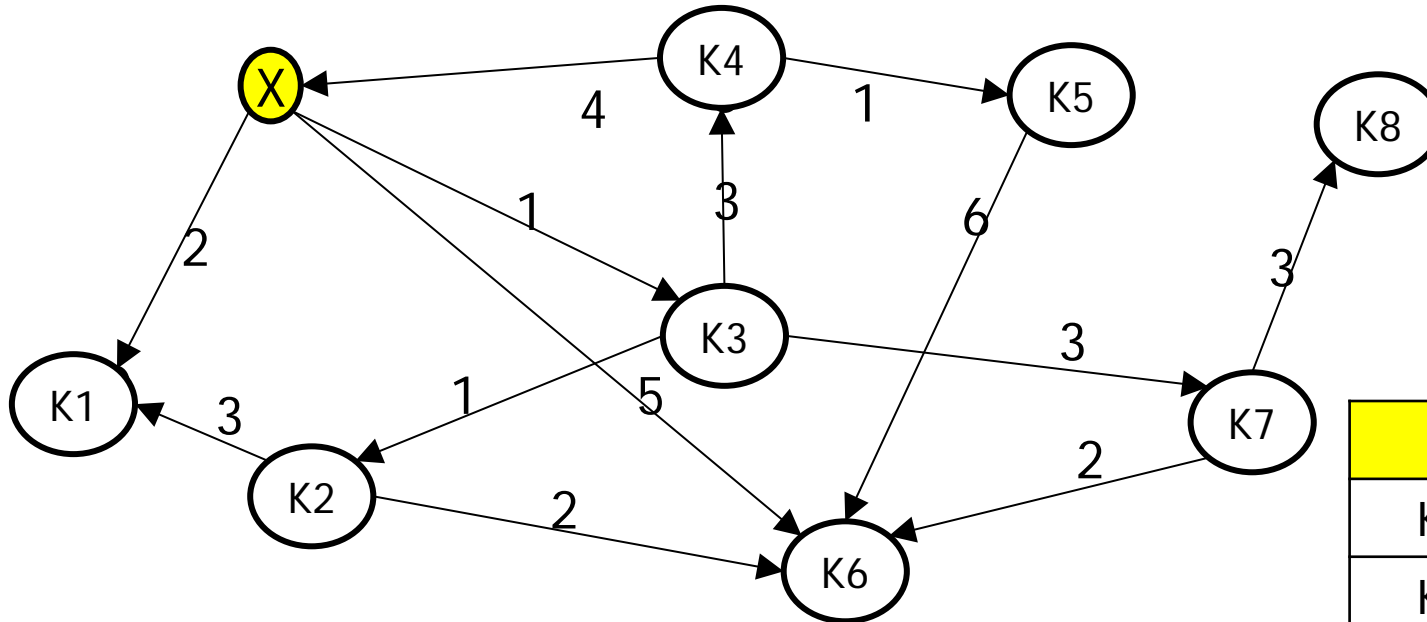
Example for Algorithm



- Pick x

| X | 0 |
|----|---|
| K1 | |
| K2 | |
| K3 | |
| K4 | |
| K5 | |
| K6 | |
| K7 | |
| K8 | |

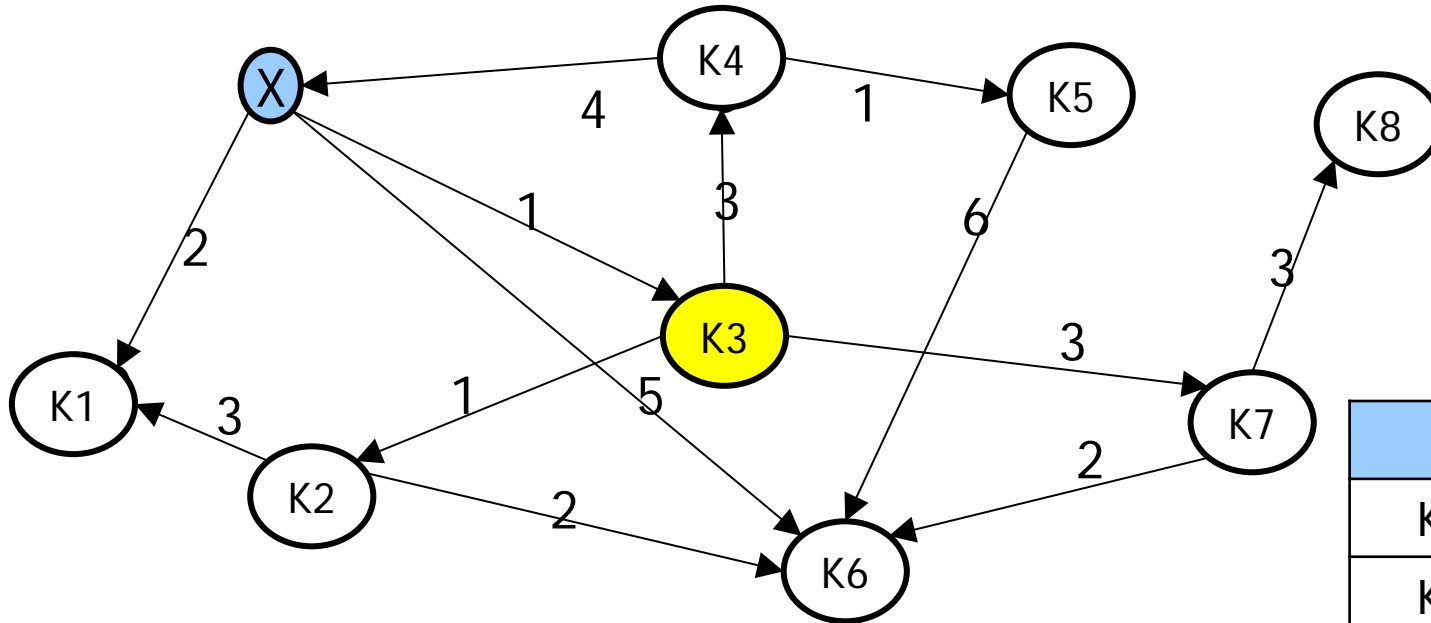
Example for Algorithm



| X | 0 |
|----|---|
| K1 | 2 |
| K2 | |
| K3 | 1 |
| K4 | |
| K5 | |
| K6 | 5 |
| K7 | |
| K8 | |

- Pick x
- Adapt distances to all neighbors

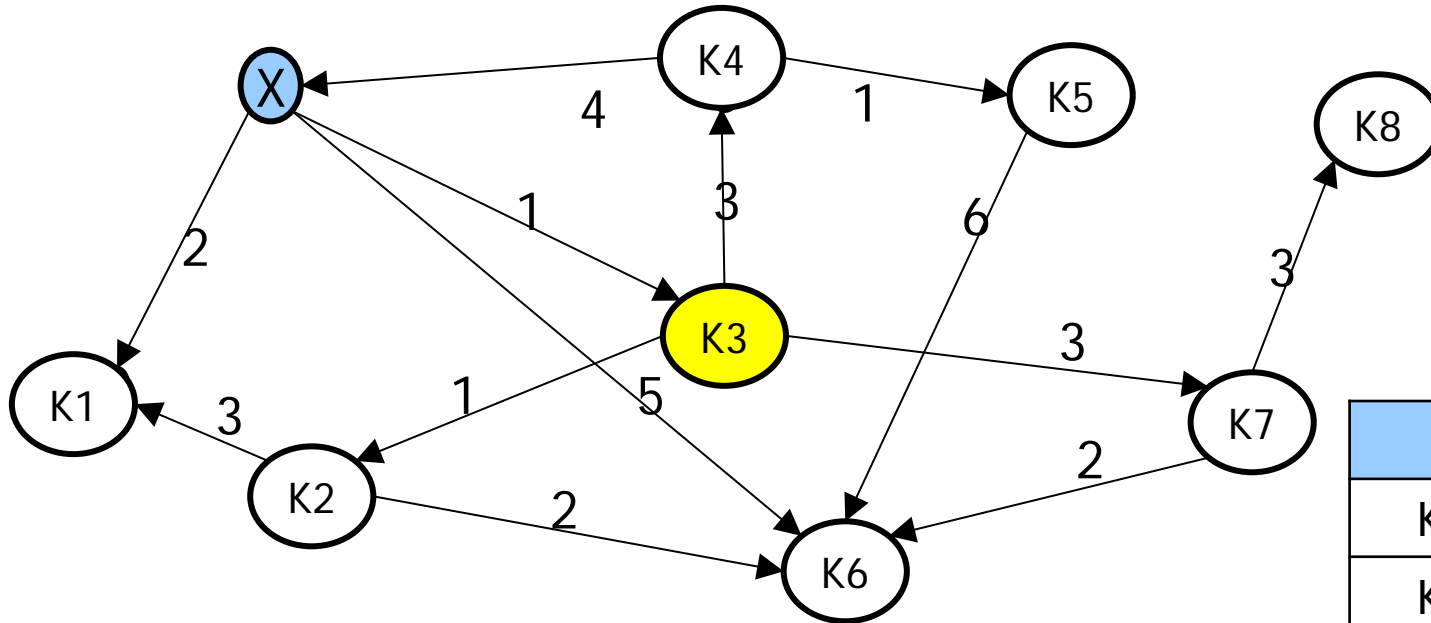
Example for Algorithm



| | |
|----|---|
| X | 0 |
| K1 | 2 |
| K2 | |
| K3 | 1 |
| K4 | |
| K5 | |
| K6 | 5 |
| K7 | |
| K8 | |

- Pick K3 (closest to x)

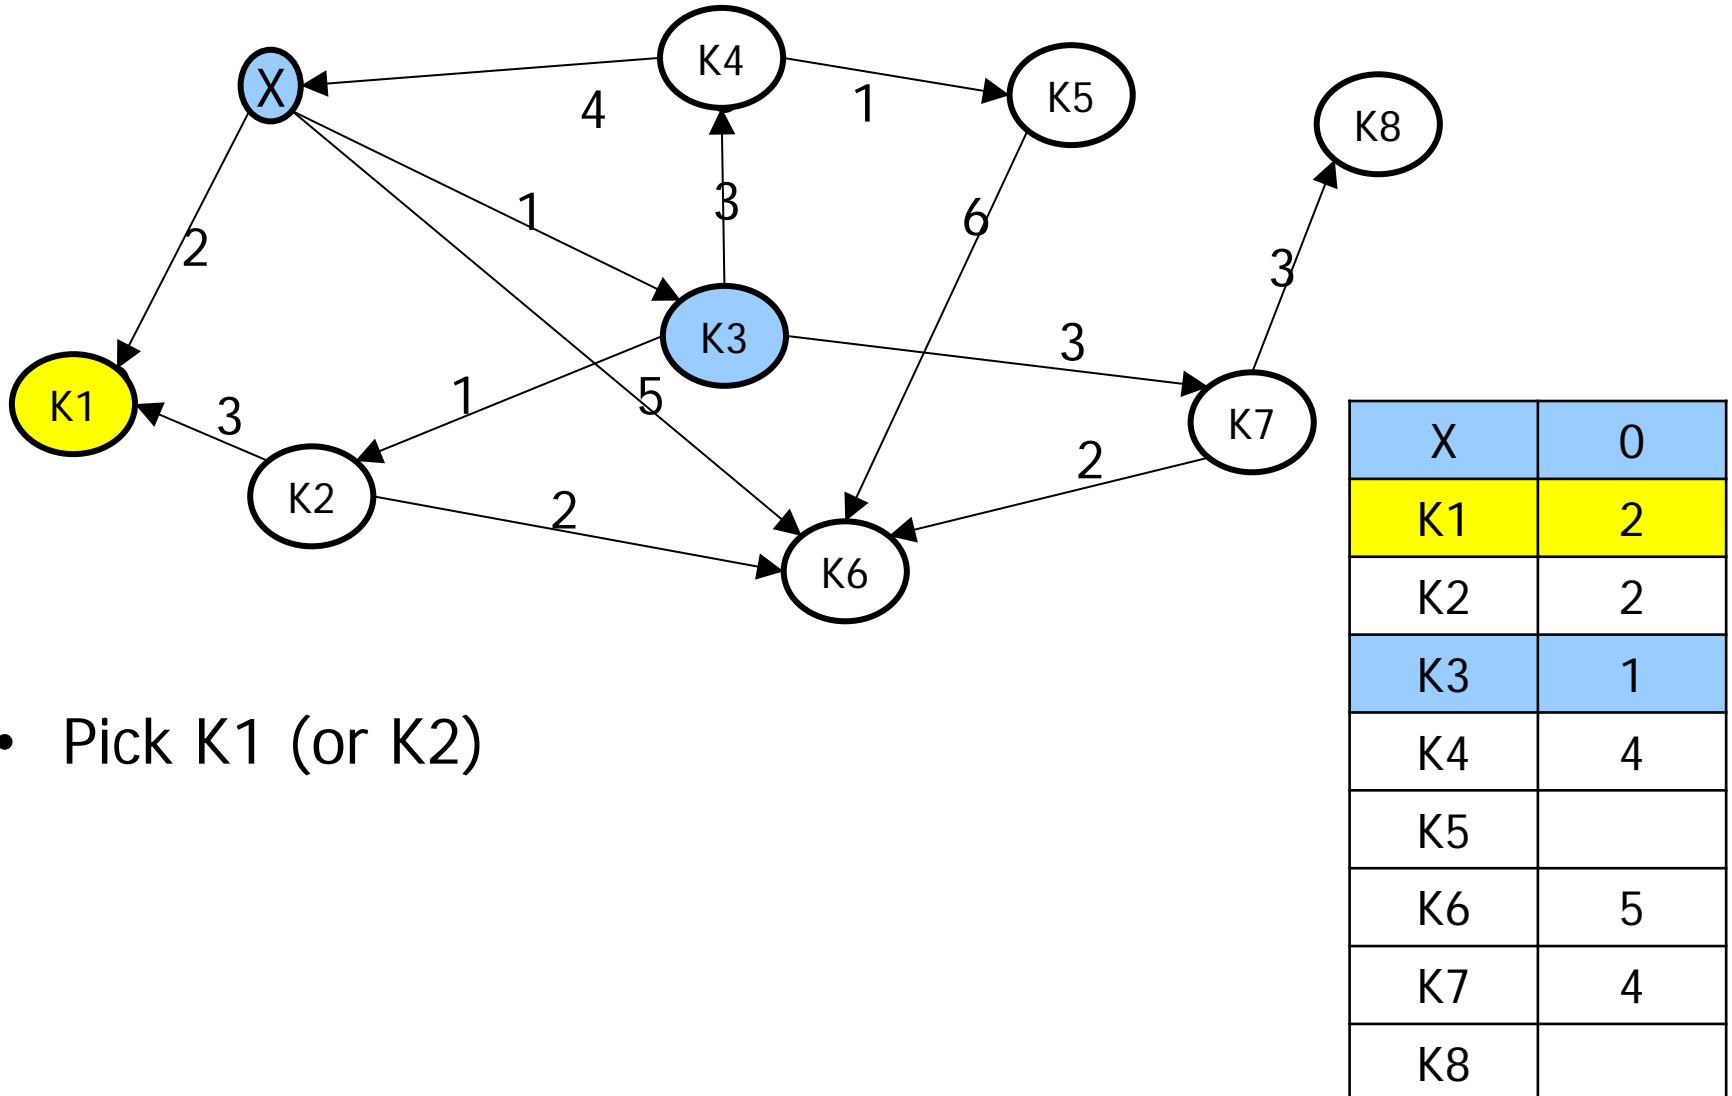
Example for Algorithm



| | |
|----|---|
| X | 0 |
| K1 | 2 |
| K2 | 2 |
| K3 | 1 |
| K4 | 4 |
| K5 | |
| K6 | 5 |
| K7 | 4 |
| K8 | |

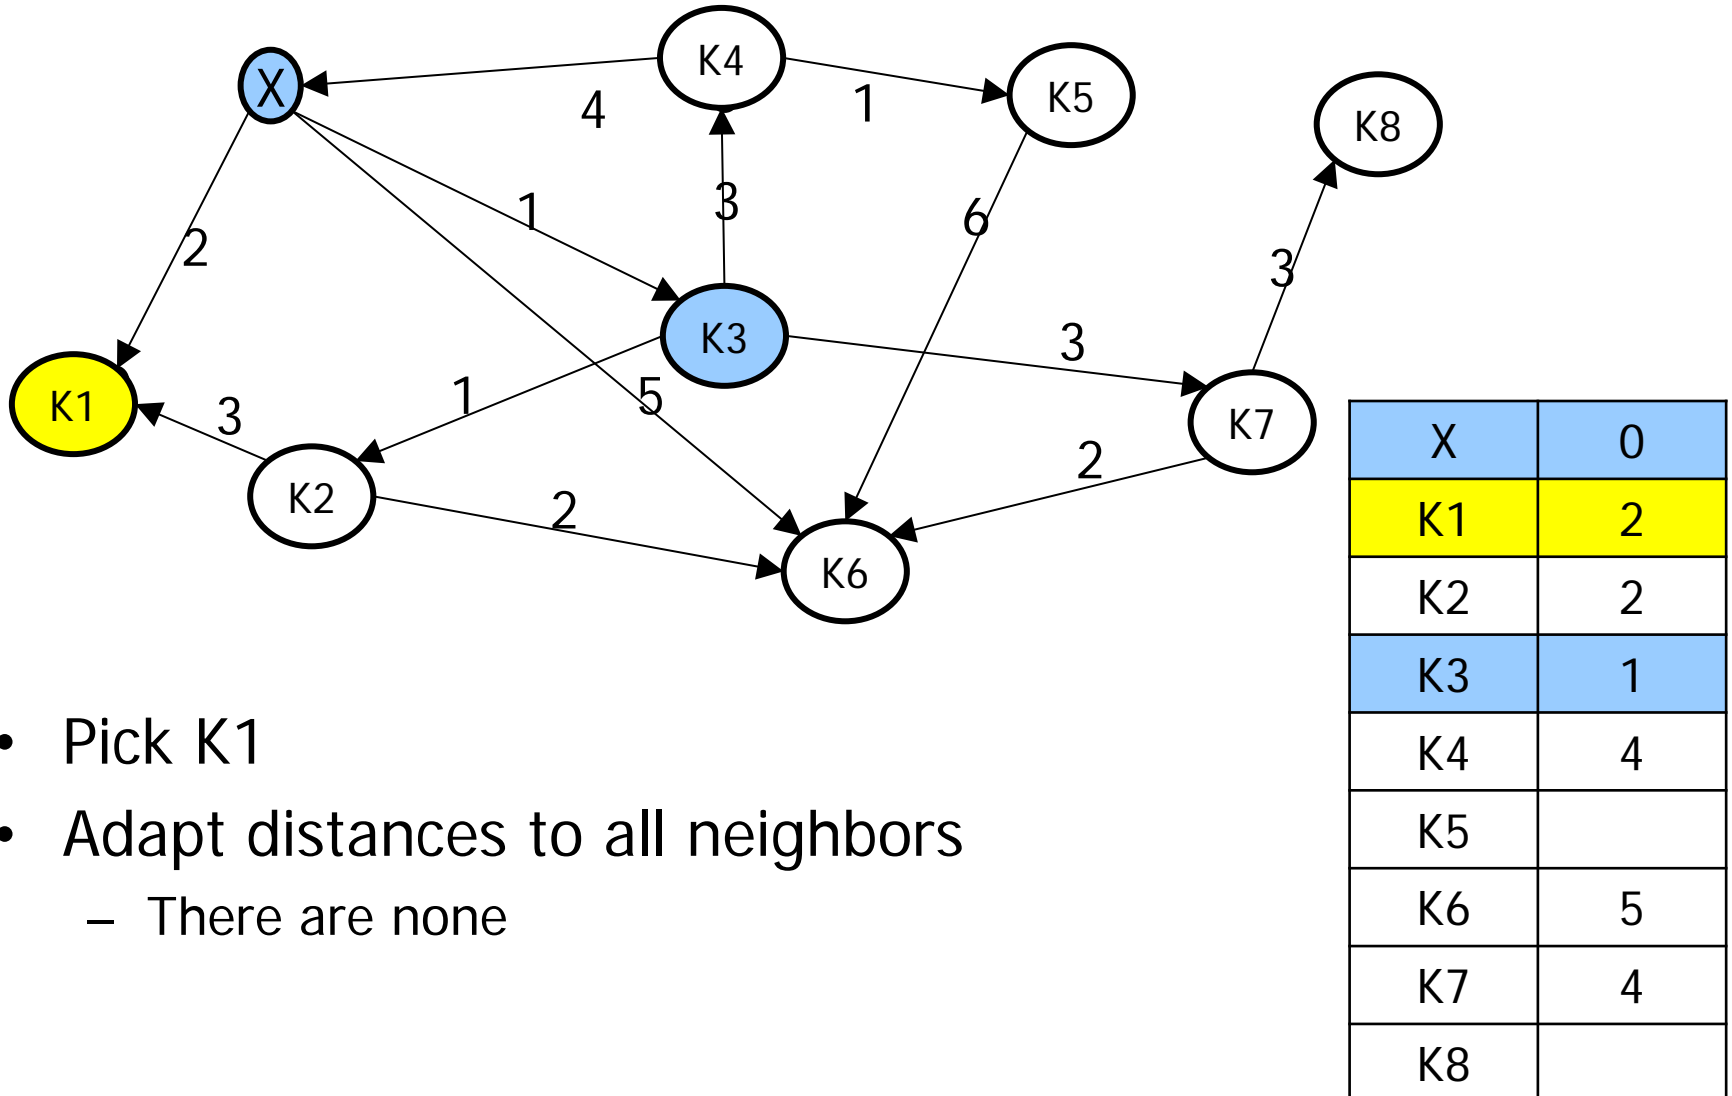
- Pick K3
- Adapt distances (from x) to all neighbors (of K3)

Example for Algorithm



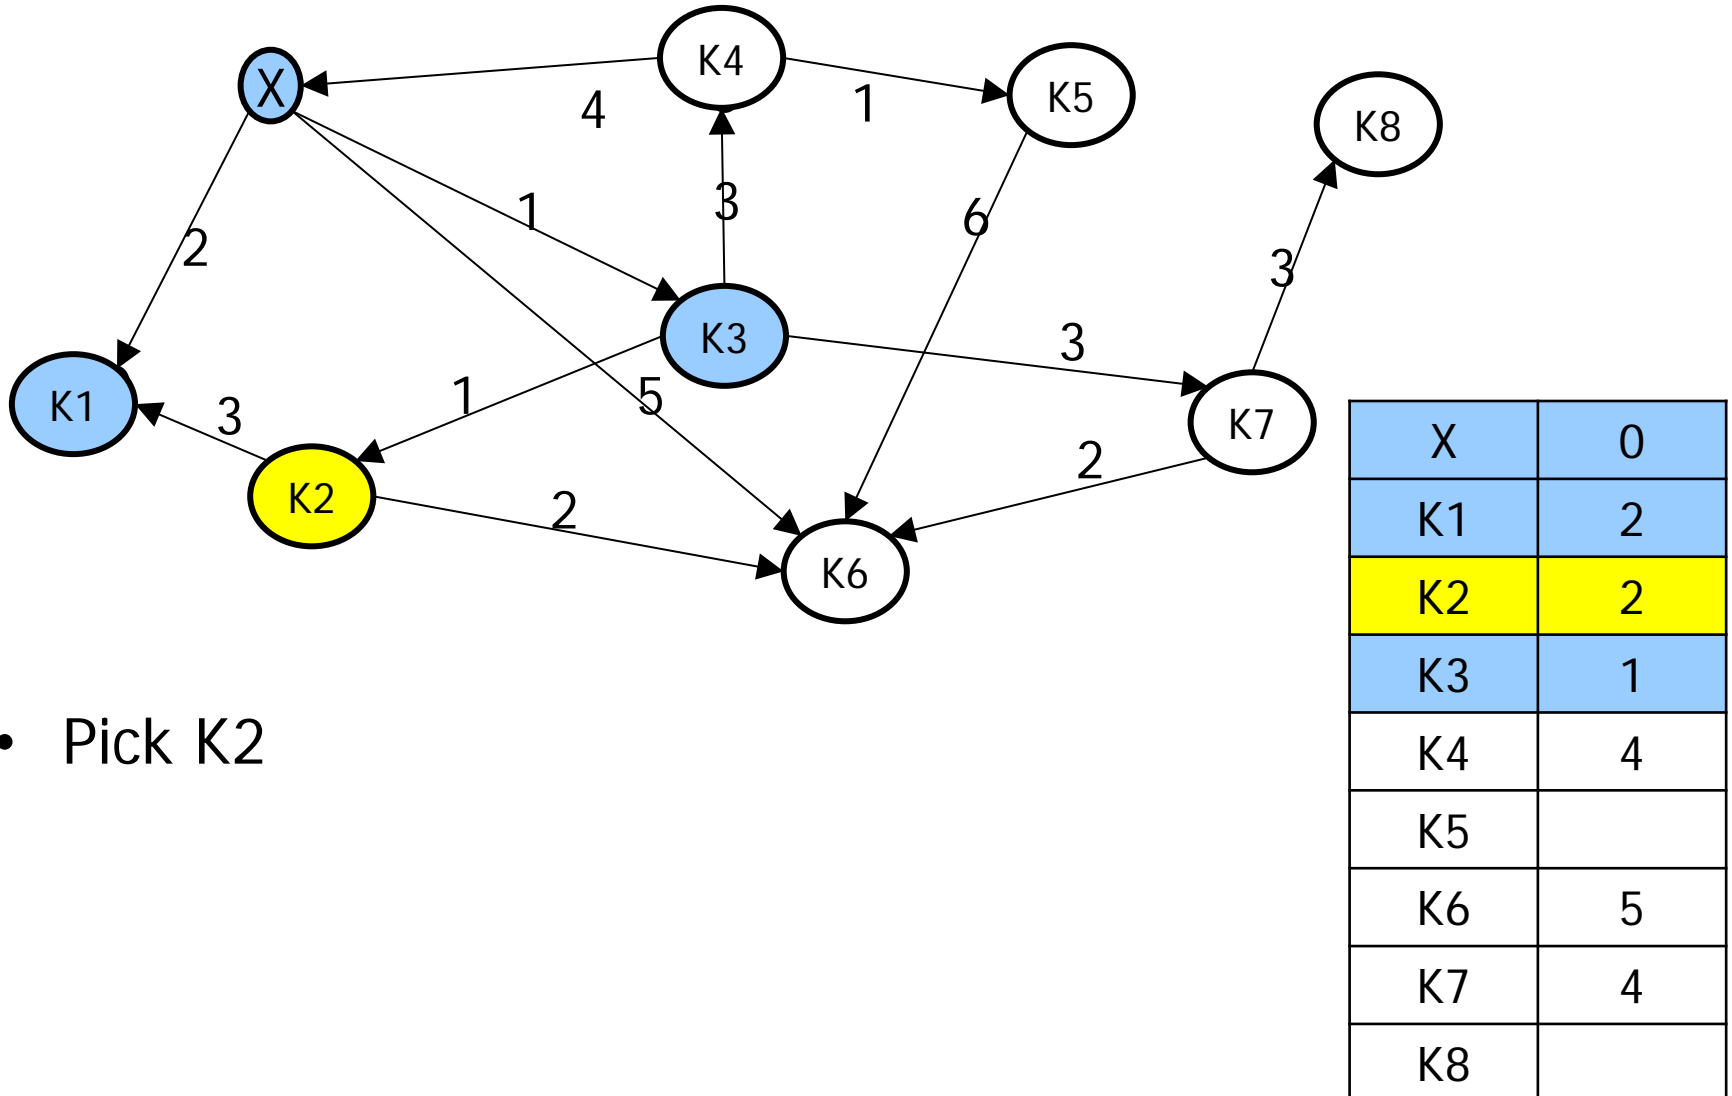
- Pick K1 (or K2)

Example for Algorithm



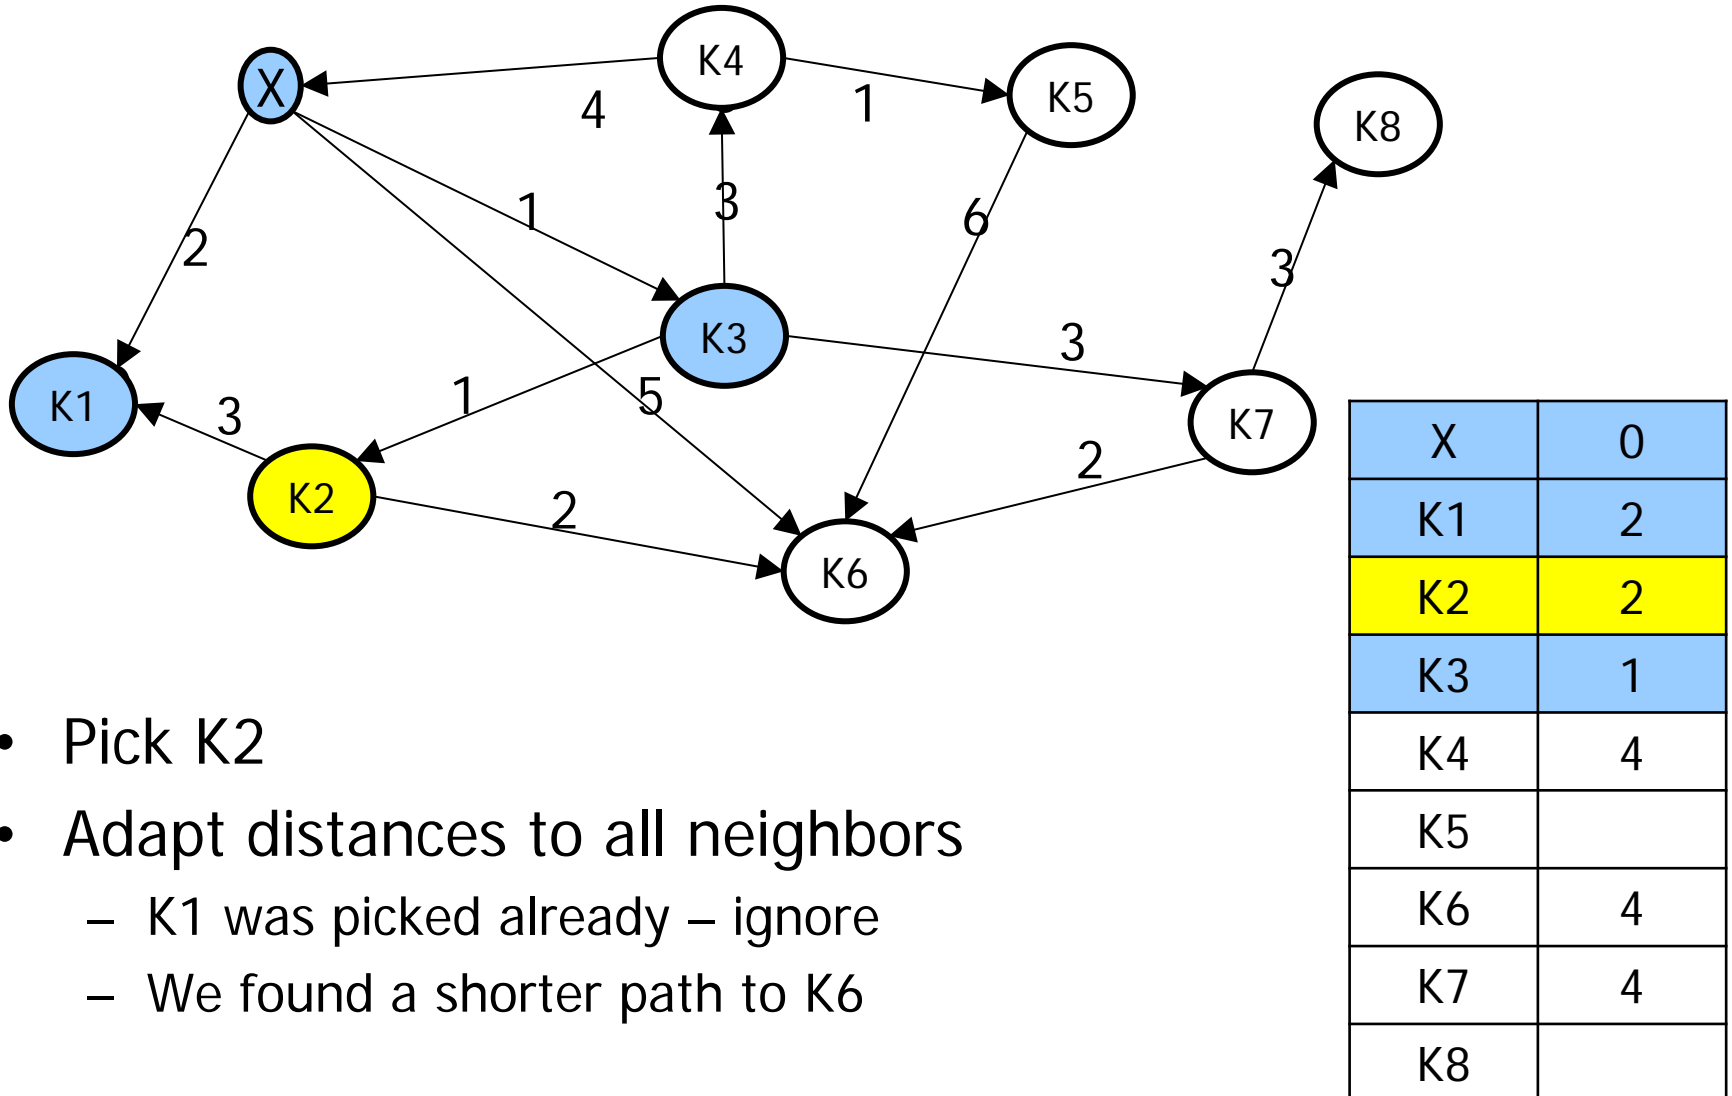
- Pick K1
- Adapt distances to all neighbors
 - There are none

Example for Algorithm



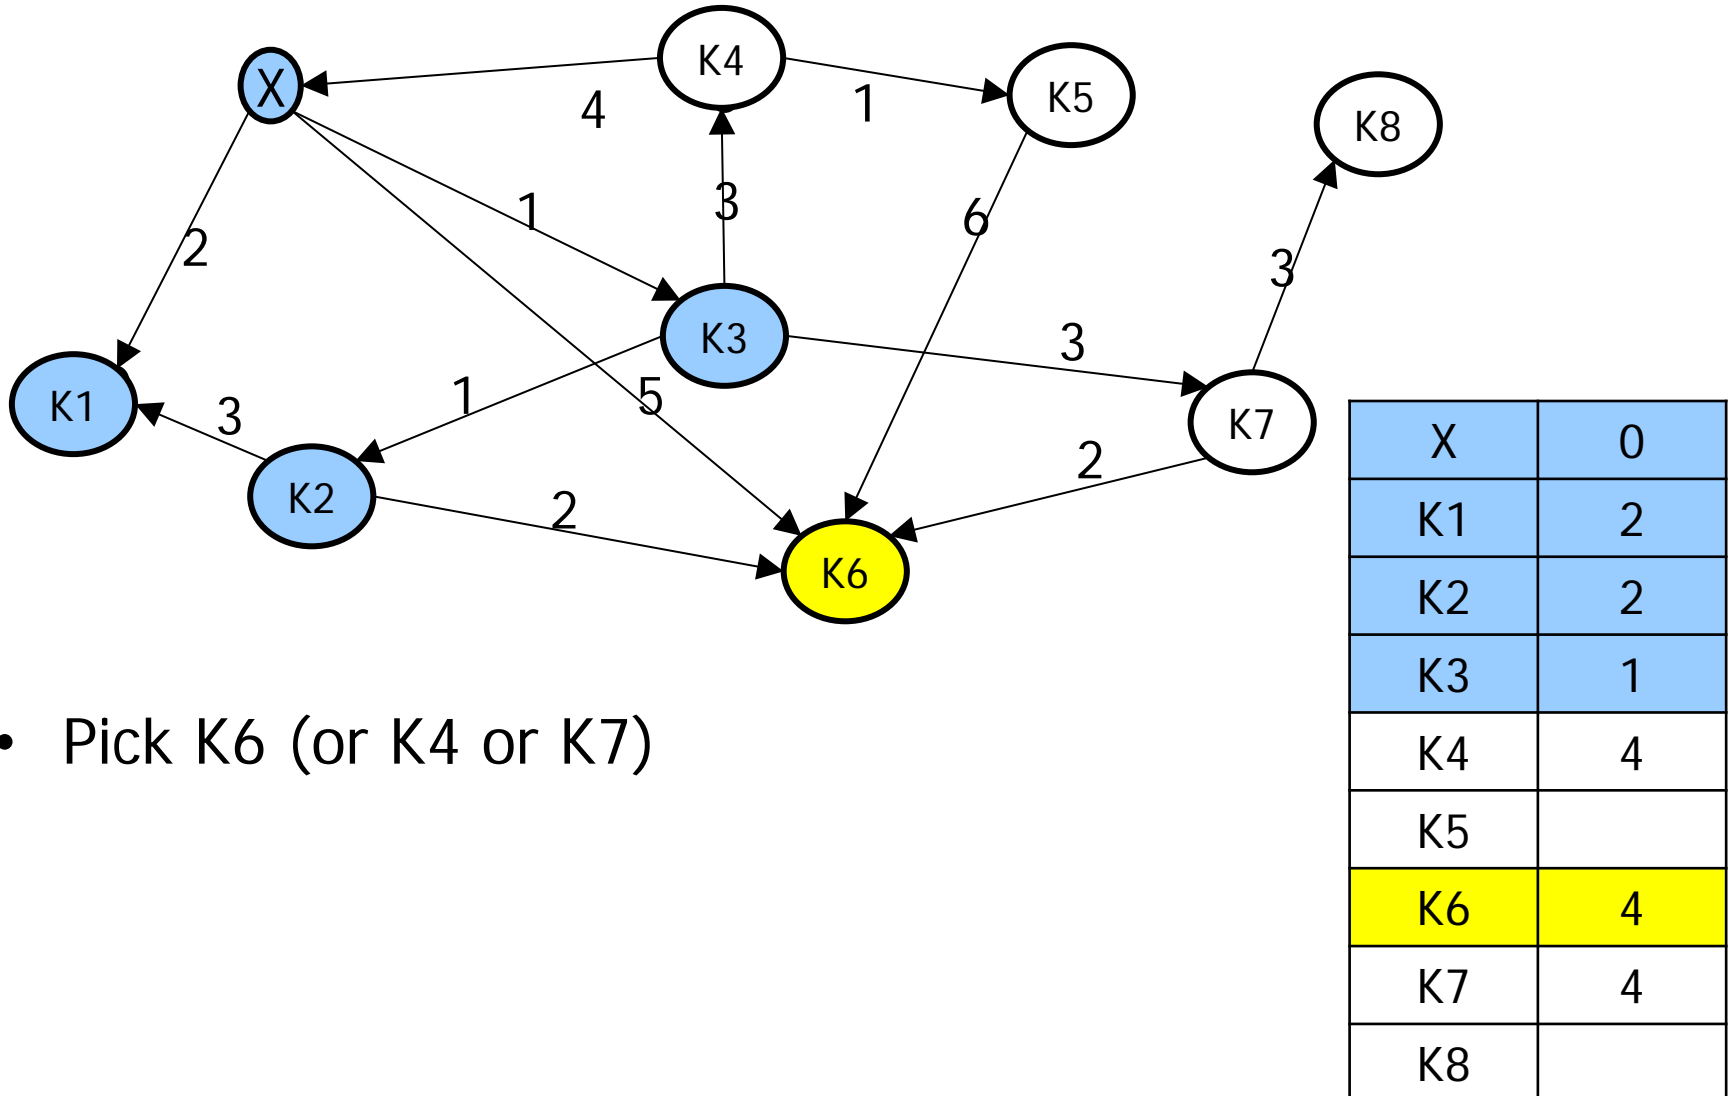
- Pick K2

Example for Algorithm

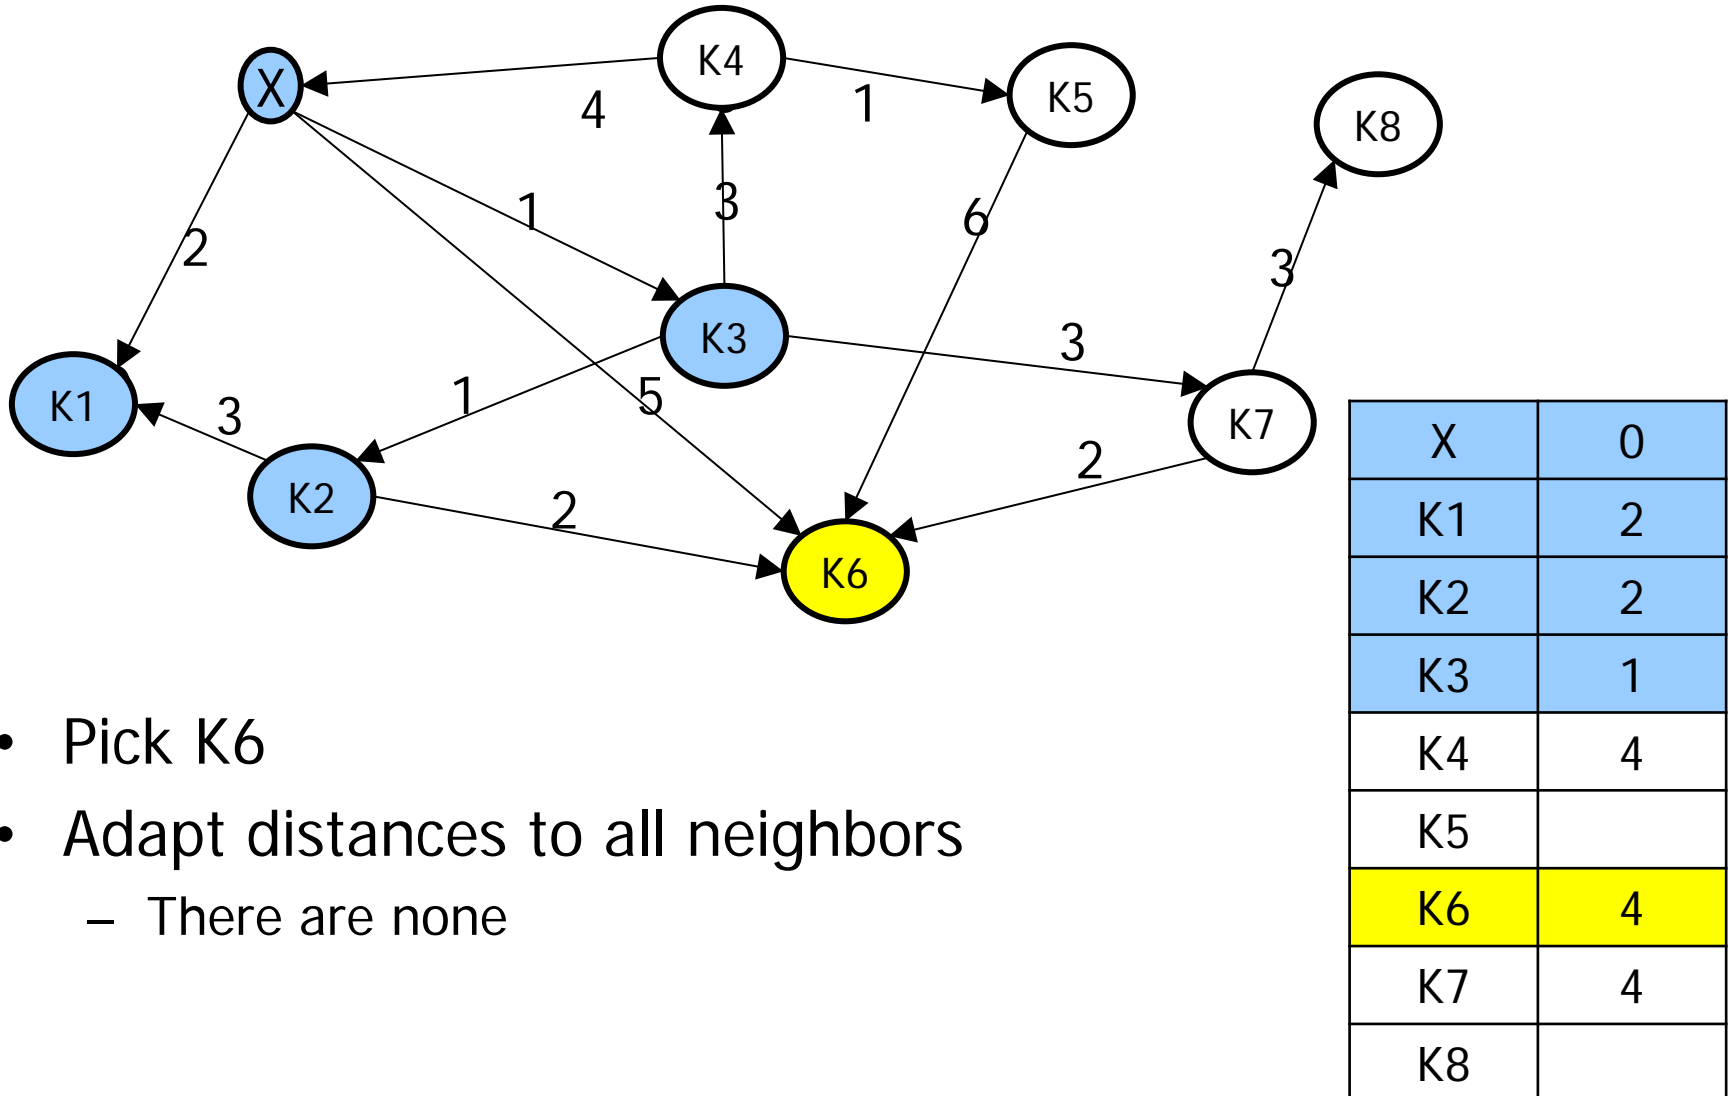


- Pick K2
- Adapt distances to all neighbors
 - K1 was picked already – ignore
 - We found a shorter path to K6

Example for Algorithm

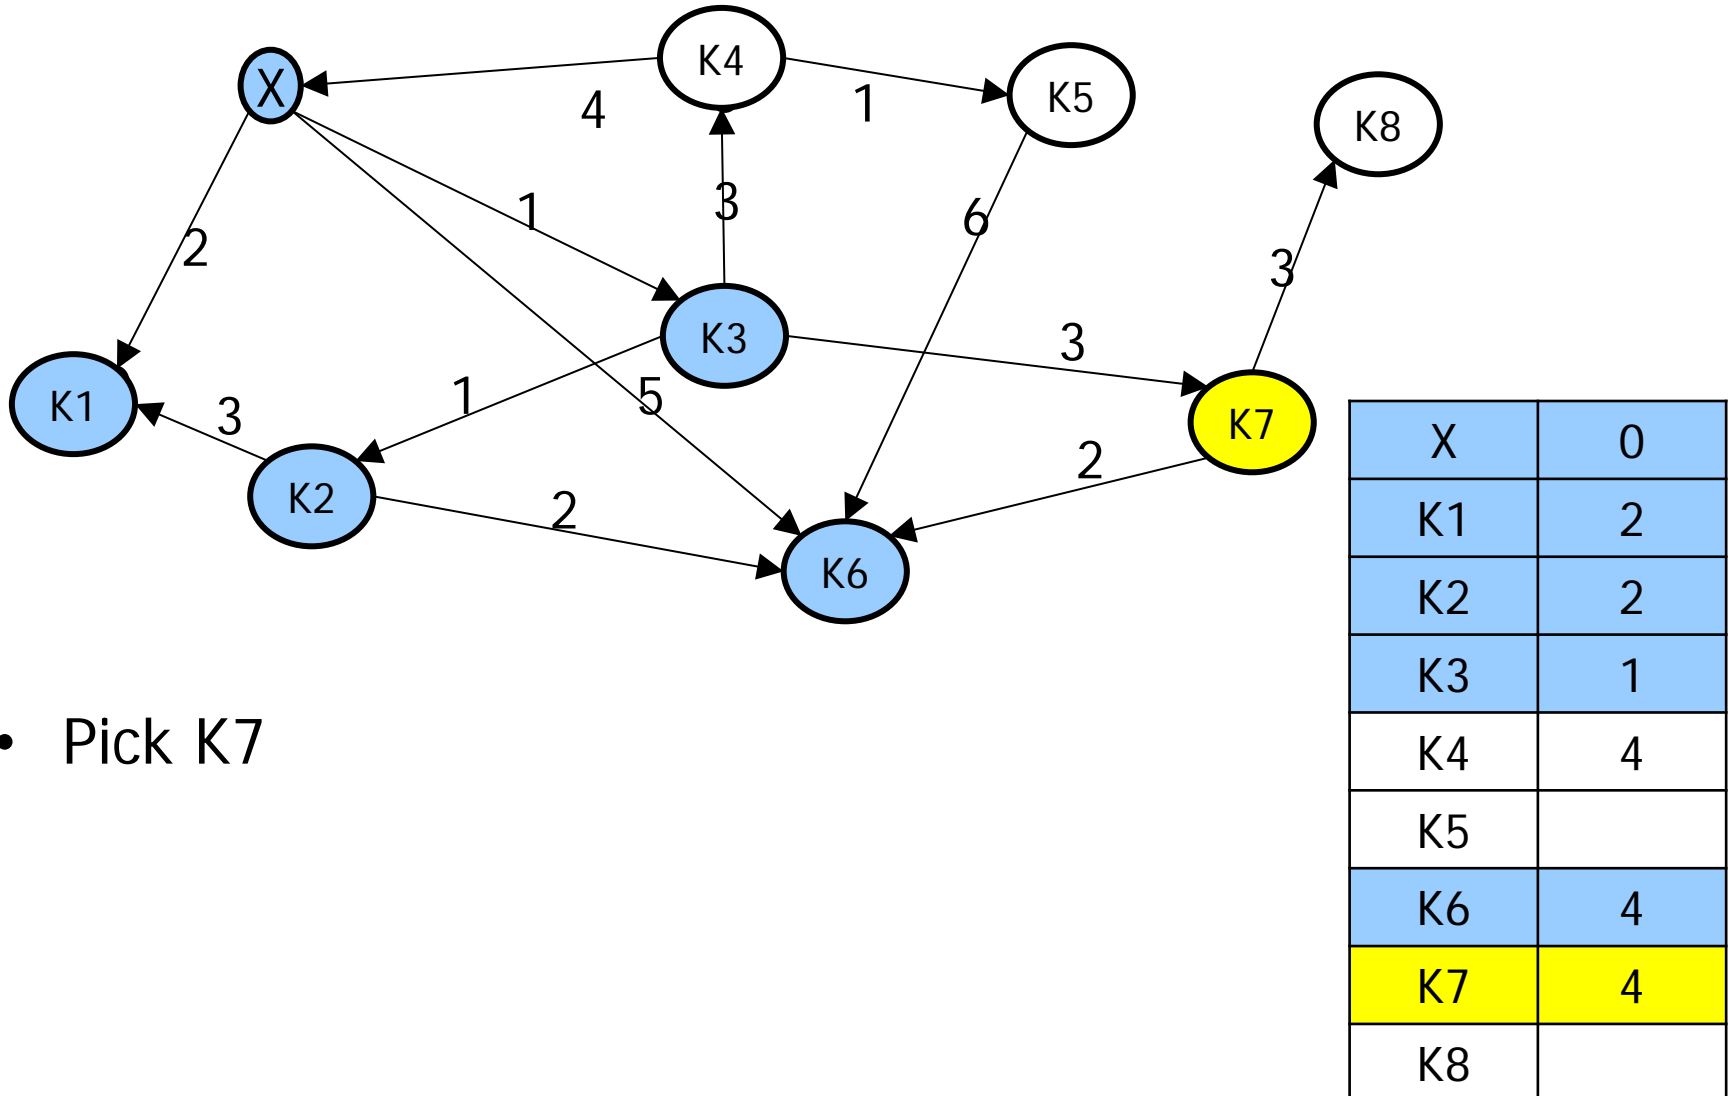


Example for Algorithm

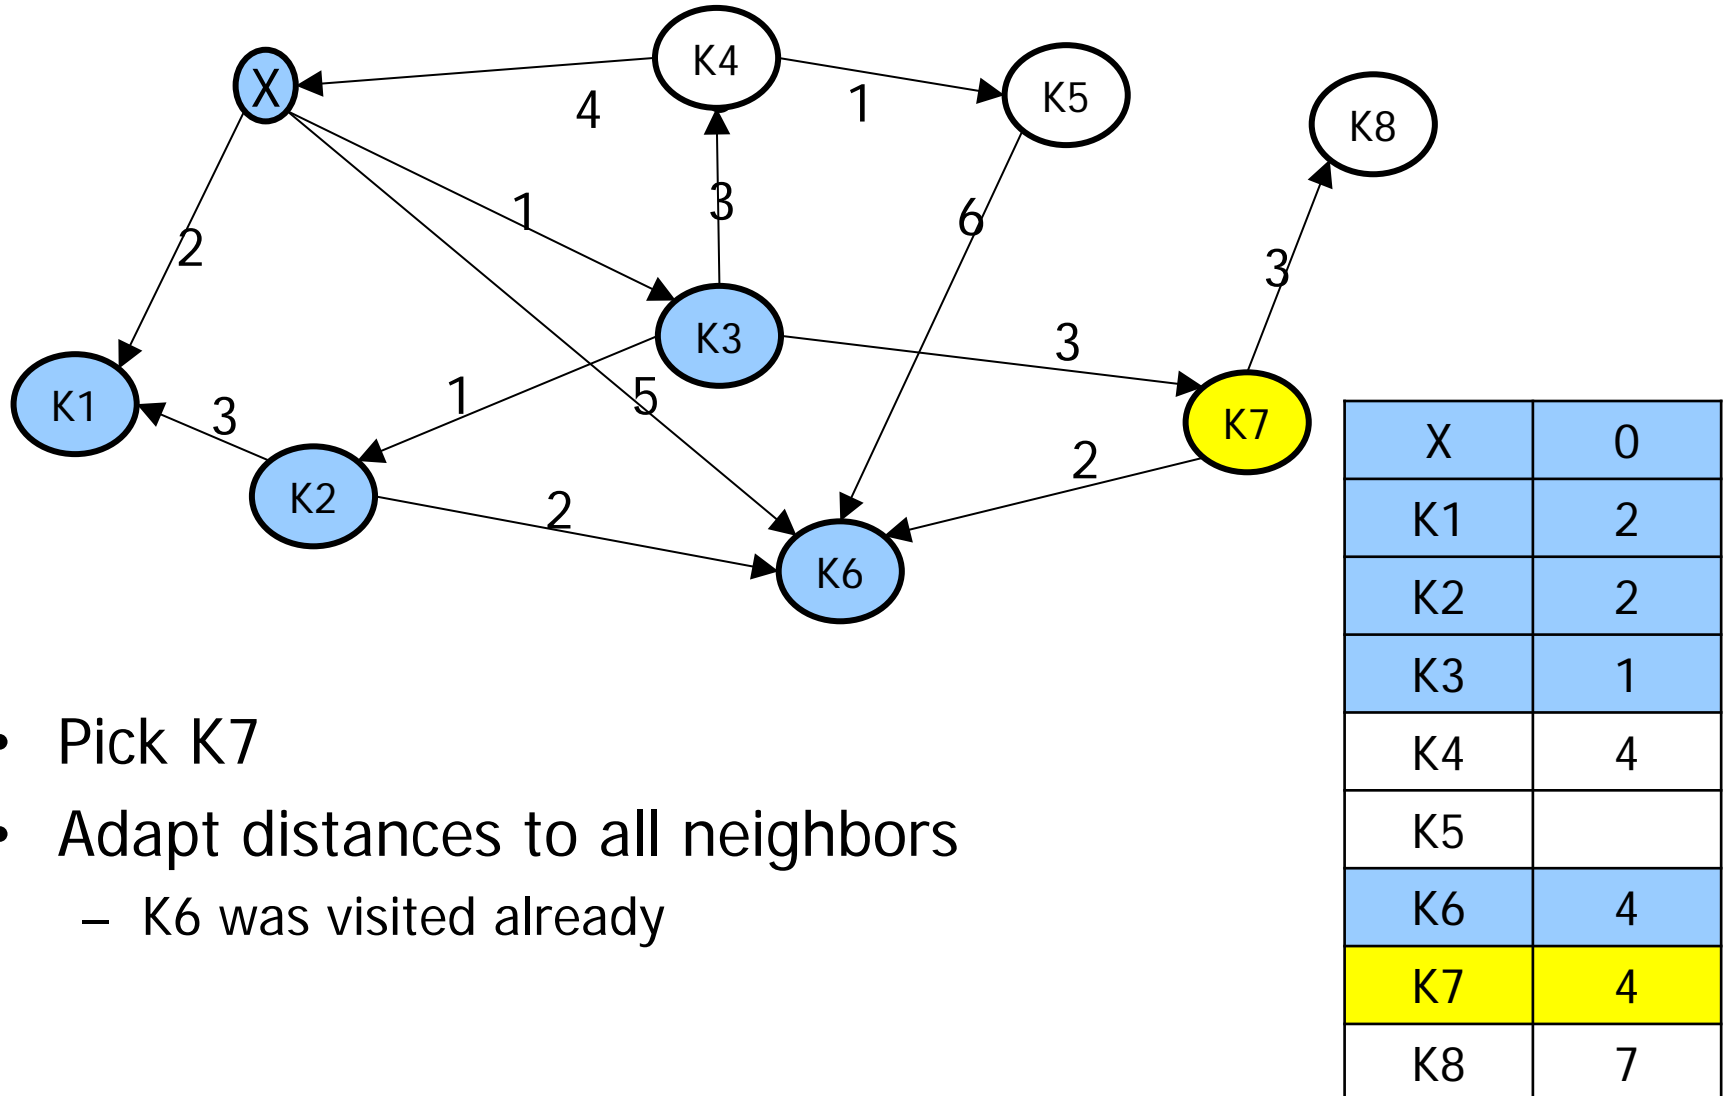


- Pick K6
- Adapt distances to all neighbors
 - There are none

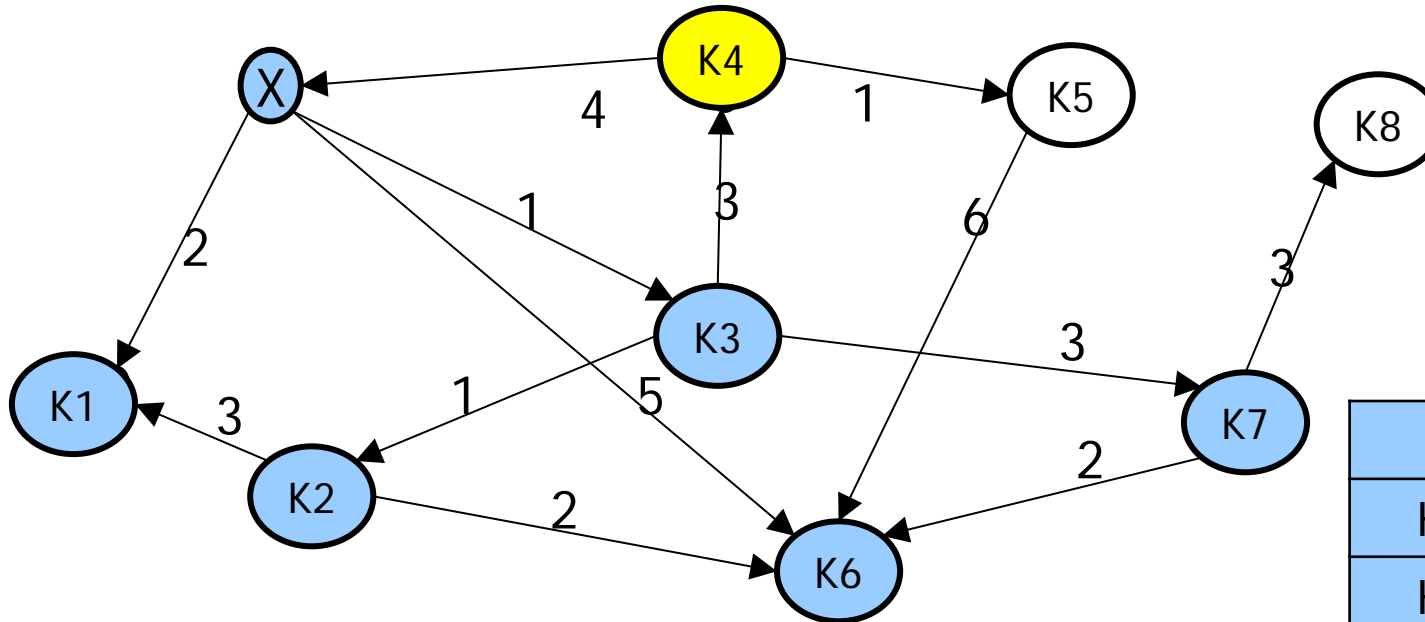
Example for Algorithm



Example for Algorithm



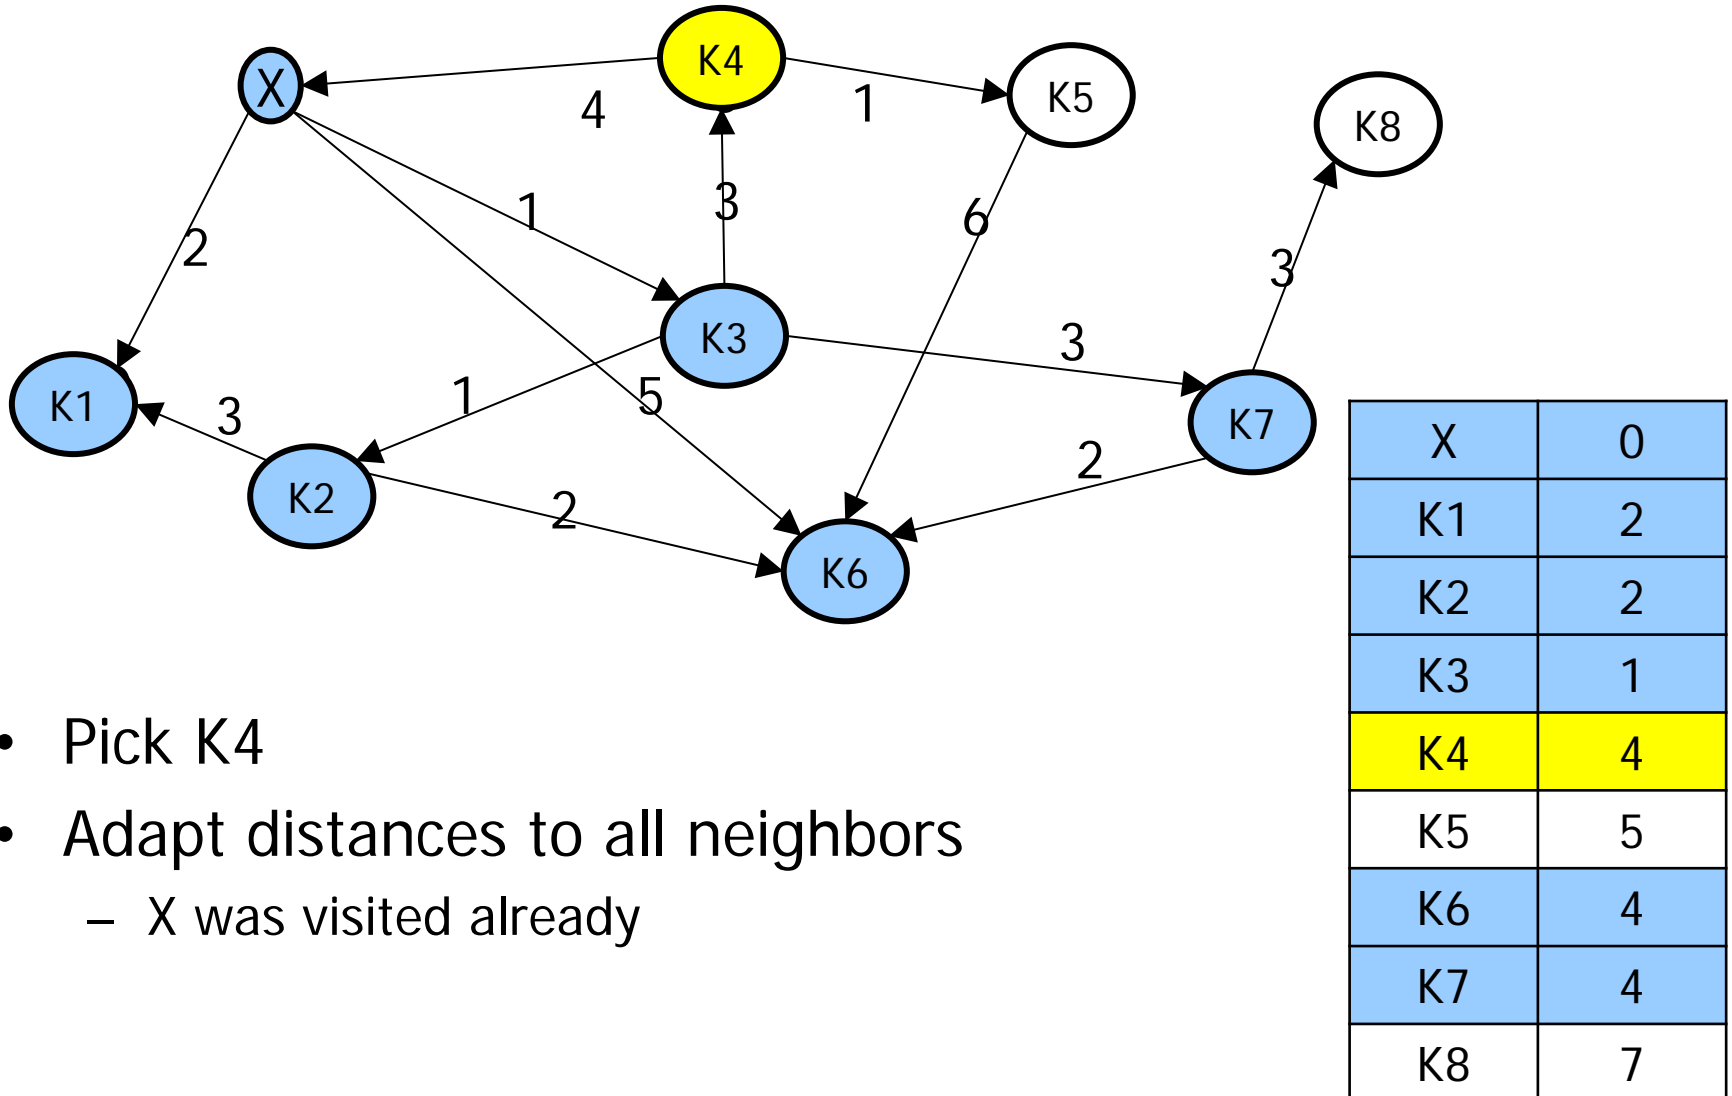
Example for Algorithm



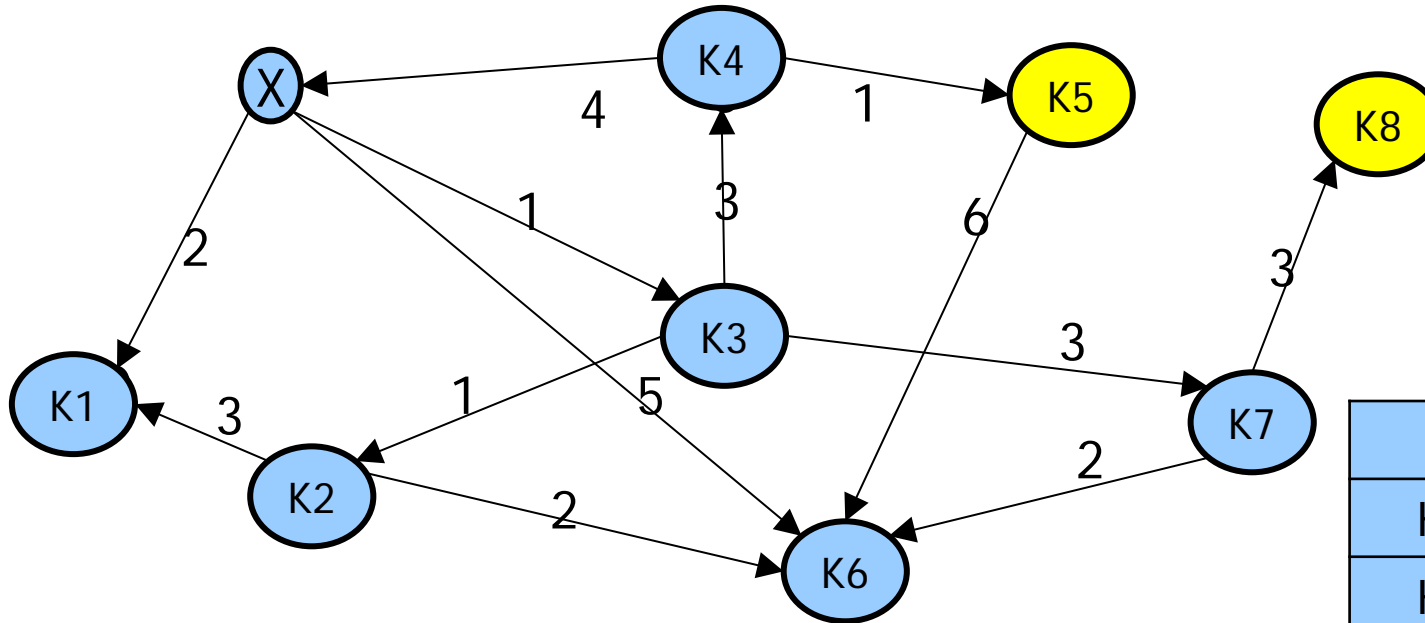
- Pick K4

| | |
|----|---|
| X | 0 |
| K1 | 2 |
| K2 | 2 |
| K3 | 1 |
| K4 | 4 |
| K5 | |
| K6 | 4 |
| K7 | 4 |
| K8 | 7 |

Example for Algorithm



Example for Algorithm



- Pick K5 ... Pick K8

| | |
|----|---|
| X | 0 |
| K1 | 2 |
| K2 | 2 |
| K3 | 1 |
| K4 | 4 |
| K5 | 5 |
| K6 | 4 |
| K7 | 4 |
| K8 | 7 |

A Closer Look

```
1. G = (V, E);
2. x : start_node;      # x ∈ V
3. A : array_of_distances;
4. ∀i: A[i] := ∞;
5. L := V;
6. A[x] := 0;
7. while L ≠ ∅
8.   k := L.get_closest_node();
9.   L := L \ k;
10.  forall (k,f,w) ∈ E do
11.    if f ∈ L then
12.      new_dist := A[k] + w;
13.      if new_dist < A[f] then
14.        A[f] := new_dist;
15.      end if;
16.    end if;
17.  end for;
18. end while;
```

- Algorithm seems to work
 - **Proof** and analysis will follow later
 - Hint: 8 is passed-by $|V|$ times and 12 at most $|E|$ times
- Central: **get_closest_node()**
 - Needs to find the node k in L for which $A[k]$ is the smallest
 - $A[k]$ is changed a lot during a run
- Searching A ? Resorting A ?
- Better: **Priority queue**
 - List of tuples (o, v) (object,value)
 - Central operation: Return tuple where v is smallest

Content of this Lecture

- Priority Queues
- Using Heaps
- Using Fibonacci Heaps

Priority Queues

- A **priority queue** (PQ) is an ADT with 3 essential operations
 - **add(o, v)**: Add element o with value (priority) v
 - May be also bulk insert – convert a list in a priority queue
 - **getMin()**: Retrieve **element with highest priority**
 - **removeMin()**: Remove element with smallest value
- Typical additional operations
 - **merge(p1, p2)**: Merge two PQs into one (properly sorted)
 - **delete(o)**: Delete o from PQ
 - **changeValue(o, v)**: Change value of o to v

Applications

- Games (e.g. chess)
 - The machine explores next movements but cannot look at all of them; give each move an assumed benefit and explore **moves with highest benefit first** (also called **A* algorithm**)
- Event simulators
 - While events are handled, new events are generated for the future; manage all events in a PQ sorted by event time and **always pull the next event**
- Quality of Service in a network
 - When bandwidth is limited, sort all transmission requests in a PQ and **transmit by highest priority**
- ...

Naive Implementations (with $|Q|=n$)

- Using a linked list
 - **add** requires $O(1)$
 - **getMin** requires $O(n)$ [bad]
 - **deleteMin** requires $O(1)$ (if we keep the pointer after a getMin)
 - **merge** requires $O(1)$
- Using a linked list sorted by priority
 - **add** requires $O(n)$ [bad]
 - **getMin** requires $O(1)$
 - **deleteMin** requires $O(1)$
 - **merge** requires $O(n+m)$

Maybe Arrays?

- Using a sorted array
 - **add** requires $O(n)$ [We find the position in $\log(n)$, but then have to free a cell by moving all elements after this cell]
 - **getMin** requires $O(1)$
 - **deleteMin** requires $O(n)$
- PQs are typically used in applications where elements are inserted and removed all the time
- We need a DS that can **change its size dynamically** at very low cost
- We want **constant or at most log-time** for all operations

Content of this Lecture

- Priority Queues
- Using Heaps
 - Heaps
 - Operations on Heaps
 - Heap Sort
- Using Fibonacci Heaps

Heap-based PQ


- Unsorted lists require $O(n)$ for `getMin()`
 - We don't know where the smallest element is
- Sorted lists require $O(n)$ for `add()`
 - We don't know where to put the new element
- Can we find a way to keep the list “a little sorted”?
 - Actually, we only want the smallest element at a fixed position
 - All other elements can be at arbitrary places
 - `add()` / `deleteMin()` should be faster than $O(n)$, because they don't need to keep the entire list sorted
- One such structure is called a **heap**

Heaps

- Definition


A *heap* is a labeled binary tree for which the following holds

- *Form-constraint* (FC): The *tree is complete* except the last level
 - I.e.: Every node has exactly two children
- *Heap-constraint* (HC): The value of any node is smaller than that of its children



Properties

- Order
 - A heap is “a little” sorted: We know the **smallest element** (root)
 - We know the **order for some pairs** of elements (parent-successors), but for many pairs we don’t know which is bigger (e.g. nodes in the same level)
- Size
 - A complete binary tree with m levels has $2^m - 1$ nodes
 - A heap with m levels thus has between $2^{m-1} + 1$ and $2^m - 1$ nodes
 - A heap with n nodes **has $\text{ceil}(\log(n+1))$ levels**

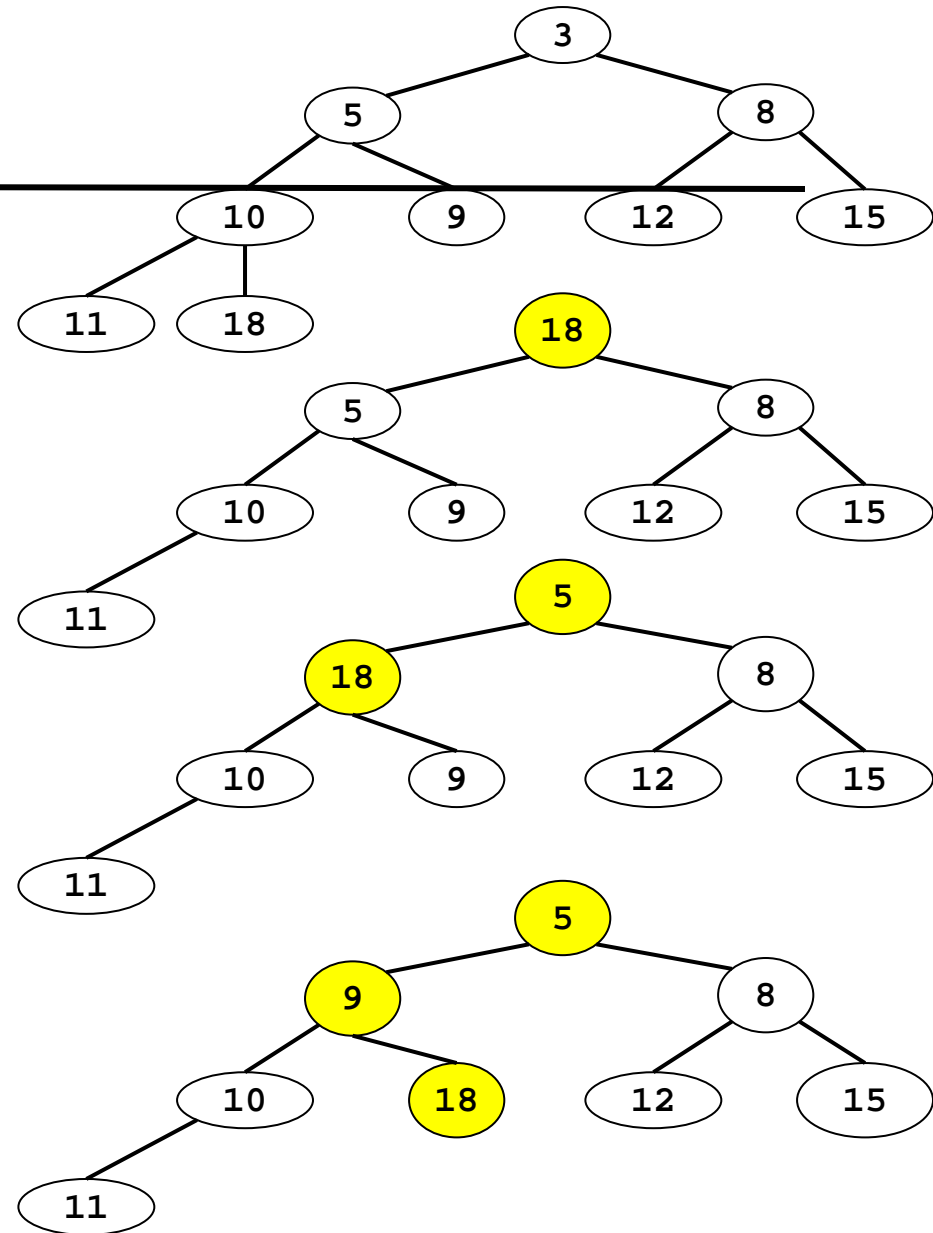


Operations

- Assume we store our PQ as a heap
- Clearly, `getMin()` is possible in $O(1)$
 - Keep a pointer to the root
- But ...
 - How can we perform `deleteMin()` – such that the new structure again is a heap?
 - How can we add an element to a heap – such that the new structure again is a heap?
 - How can we turn a list into a heap?

DeleteMin()

- We first remove the root
 - Creates **two heaps**
 - We must connect them again
- We take the „last“ node, place it in root, and **sift it down the tree**
 - Last node: right-most in the last level (actually, we can take any from the last level)
 - **Sifting down**: Exchange with smaller of both children as long as at least one child is smaller than the node itself



Analysis - Correctness





- We need to show that **FC and HC still hold**
- HC: Look at the tree after we moved a node k . k may
 - ... be smaller than its children. Then HC holds and we are done
 - ... be larger than at least one child k_2 . Assume that k_2 is the smaller of the two children (k_1, k_2) of k . We next swap k and k_2 . The **new parent (k_2) now is smaller** than its children (k_1, k), so the HC holds
 - After the last swap, k has no children – HC holds
- FC: We remove one node, then we sift down
 - Removing last node doesn't affect FC as we remove in the last level
 - Sifting does not change the **topology of the tree** (we only swap)

Analysis - Complexity

- Recall that a heap with n nodes has $\text{ceil}(\log(n+1))$ levels
- During sifting, we perform at most one comparison and one swap **in every level**
- Thus: $O(\text{ceil}(\log(n+1))) = O(\log(n))$

Add() on a Heap

- Cannot simply add on top
- Idea: We add new element somewhere in last level and **sift up**
 - We might need a new level
 - Sifting up: Compare to parent and swap **if parent is larger**



Analysis



- Correctness
 - HC
 - If parent has **only one child**, HC holds after each swap
 - Assume a parent k has children k_1 and k_2 , k_2 was swapped there in the last move, and $k_2 < k$. Since HC held before, $k < k_1$, **thus $k_2 < k < k_1$** . We swap k_2 and k , and thus the **new parent is smaller** than its children. On the other hand, if $k_2 \geq k$, HC holds immediately (and we don't swap).
 - FC: See `deleteMin()`
- Complexity: $O(\log(n))$
 - See `deleteMin()`

How to Find the Next Free / Last Occupied Node

- What do we need to find?
 - For **deleteMin**, we use the right-most leaf on the last level
 - For **add**, we add after the leaf right from the last leaf
- We actually need the parent k
 - From n , we can compute in $O(1)$ the position p of the last leaf in the last level: $p = n - 2^{\lfloor \log(n) \rfloor}$
 - Or $\log(n+1)$ for **add**
 - The parent k of the node at p has index $\lfloor p/2 \rfloor$ 'th in level $d-1$
 - The parent k' of k has index $\lfloor p/4 \rfloor$ 'th in level $d-2$
 - ...
 - Now, in each node, we can decide whether to go left or right
 - Fast trick: Use the binary representation of p

Illustration

- For **deleteMin**, we need x (or x'); for **add**, we need y (or y')
 - $p(x)=0, p(y)=1, p(x')=4, p(y')=5$
 - Binary: 000, 001, 100, 101
- Go through bitstring from left-to-right
- Next bit=0: Go left
- Next bit=1: Go right
- Allows finding k in $O(\log(n))$



Summary

| | Linked list | Sorted linked list | Heap |
|-------------|----------------|--------------------|--------------------------|
| getMin() | $O(n)$ | $O(1)$ | $O(1)$ |
| deleteMin() | $O(1)$ | $O(1)$ | $O(\log(n))$ |
| add() | $O(1)$ | $O(n)$ | $O(\log(n))$ |
| merge() | $O(1)$ | $O(n_1+n_2)$ | $O(\log(n_1)*\log(n_2))$ |
| Space | n add. pointer | n add. pointer | n add. pointer |

Heaps can also be kept efficiently in an array – no extra space, but limit to heap size

But merge() requires $O(n_1+n_2)$ or $O(n_1*\log(n_2+n_1))$ when using an array

Creating a Heap

- We start with an unsorted list with n elements
- Naïve algorithm: Start with empty heap and perform n additions
 - Obviously requires $O(n \cdot \log(n))$
- Better: **Bottom-Up-Sift-Down**
 - Build a tree from the n elements fulfilling the FC (but not HC)
 - Simple fill a tree level-by-level – this is in $O(n)$
 - Sift-down all nodes on the second-last level
 - Sift-down all nodes on the third-last level
 - ...
 - Sift down root

Analysis

- Correctness

- After finishing one level, all subtrees starting in this level are heaps because sifting-down ensures FC and HC (see `deleteMin()`)
- Thus, when we are done with the first level (root), we have a heap

- Analysis

- We look at the cost per level h ($1 \dots \log(n)=d$)
- For every node at level h , we need at most $d-h$ operations
- At level $h \neq d$, there are 2^{h-1} nodes
 - For nodes at level d , we don't do anything
- Over all levels, this yields

$$T(n) = \sum_{h=1}^{d-1} 2^{h-1} * (d - h) = \sum_{h=1}^{d-1} h * 2^{d-h-1} = 2^{d-1} \sum_{h=1}^{d-1} \frac{h}{2^h} \leq n * \sum_{h=1}^{\infty} \frac{h}{2^h} = n * 2 = O(n)$$

Heap Sort


- Heaps also are a suitable data structure for sorting
- **Heap-Sort** (a classical one)
 - Given an unsorted list, first create a heap in $O(n)$
 - Repeat
 - Take the smallest element and store in array in $O(1)$
 - Re-build heap in $O(\log(n))$
 - Call `deleteMin(root)`
 - Until heap is empty – after n iterations
- Thus: $O(n \cdot \log(n))$
 - Average-case only slightly better
- Can be **implemented in-place** when heap is stored in array
 - See [OW93] for details

Content of this Lecture

- Priority Queues
- Using Heaps
- Using Fibonacci Heaps

Fibonacci-Heaps (very rough sketch)

- A **Fibonacci Heap (FH)** is a forest of (non-binary) heaps with disjoint values
 - All roots are maintained in a double-linked list
 - Special pointer (**min**) to the **smallest root**
 - Accessing this value (**getMin()**) obviously is $O(1)$



Source: S.Albers,
Alg&DS, SoSe 2010

Maintenance of a FH

- FHs are maintained in a **lazy fashion**
 - **add(v)**: We create a new heap with a single element node with value v. Add this **heap to the list of heaps**; adapt min-pointer, if v is smaller than previous min
 - Clearly $O(1)$
 - **merge()**: Simple **link the two root-lists** and determine new min (as min of two mins)
 - Clearly $O(1)$
- **Deleting an element** (**deleteMin()**) needs more work
 - Until now, we just added single-element heaps
 - Thus, our structure after n add() is an **unsorted list of n elements**
 - Finding the next min element after **deleteMin()** in a naïve manner would require $O(n)$

deleteMin() on FH

- Method is not complicated
 - We first remove the min element
 - We then go through the root-list and **merge heaps with the same rank** (= # of children) until all heaps in the list have different ranks
 - Merging two heaps in $O(1)$: (1) Find the heap with the smaller root value; (2) Add it as **child to the root of the other heap**
- But analysis is fairly complicated
 - The above method is $O(n)$ in worst case
 - But after every clean-up, the root-list is much smaller than before
 - Subsequent clean-ups need much less time
 - **Amortized analysis** shows: Average-case complexity is $O(\log(n))$
 - Analysis depends on the growth of the trees during merge – these grow as the **Fibonacci numbers**

Disadvantage

- Though faster on average, Fibonacci Heaps have **unpredictable delays**
- No $\log(n)$ upper bound for **every operation**
- Not suitable for real-time applications etc.

Summary

| | Linked list | Sorted linked list | Heap | Fibonacci Heap |
|-------------|--------------------|---------------------------|--------------|-----------------------|
| getMin() | $O(n)$ | $O(1)$ | $O(1)$ | $O(1)$ |
| deleteMin() | $O(1)$ | $O(n)$ | $O(\log(n))$ | $O(\log(n))^*$ |
| add() | $O(1)$ | $O(n)$ | $O(\log(n))$ | $O(1)$ |
| merge() | $O(1)$ | $O(n_1+n_2)$ | $O(\log(n))$ | $O(1)$ |