Algorithms and Data Structures

Graphs: Introduction

Ulf Leser

This Course

• Introduction	2
• Abstract Data Types	1
• Complexity analysis	1
• Styles of algorithms	1
• Lists, stacks, queues	2
• Sorting (lists)	3
• Searching (in lists, PQs, SOL)	5
• Hashing (to manage lists)	2
• Trees (to manage lists)	4
• Graphs (no lists!)	5
• Sum	**21/26**

Content of this Lecture

- Graphs
- Representing Graphs
- Traversing Graphs
- Connected Components
- Shortest Paths

Graphs

- Directed trees represent hierarchical relations
  - A directed edge can represent all kinds of relation, as long as it is
    - Asymmetric: parent_of, subclass_of, smaller_than, owns? ...
    - Cycle-free
    - Binary
- This excludes many real-life relations
  - friend_of, similar_to, reachable_by, html_linked_to, ...
- Graphs can represent all binary relationships
  - Symmetric: Undirected graphs, asymmetric: Directed graphs
- N-ary relationships: Hypergraphs
  - exam(student, professor, subject), borrow(student, book, library)

Importance

- Most graphs you will see are binary
- Most graphs you will see are simple
  - Simple graphs: At most one edge between any two nodes
  - Contrary: multigraphs
- Some graphs you will see are undirected, some directed
- Here: Only (un-)directed, binary, simple, finite graphs

Web Graph

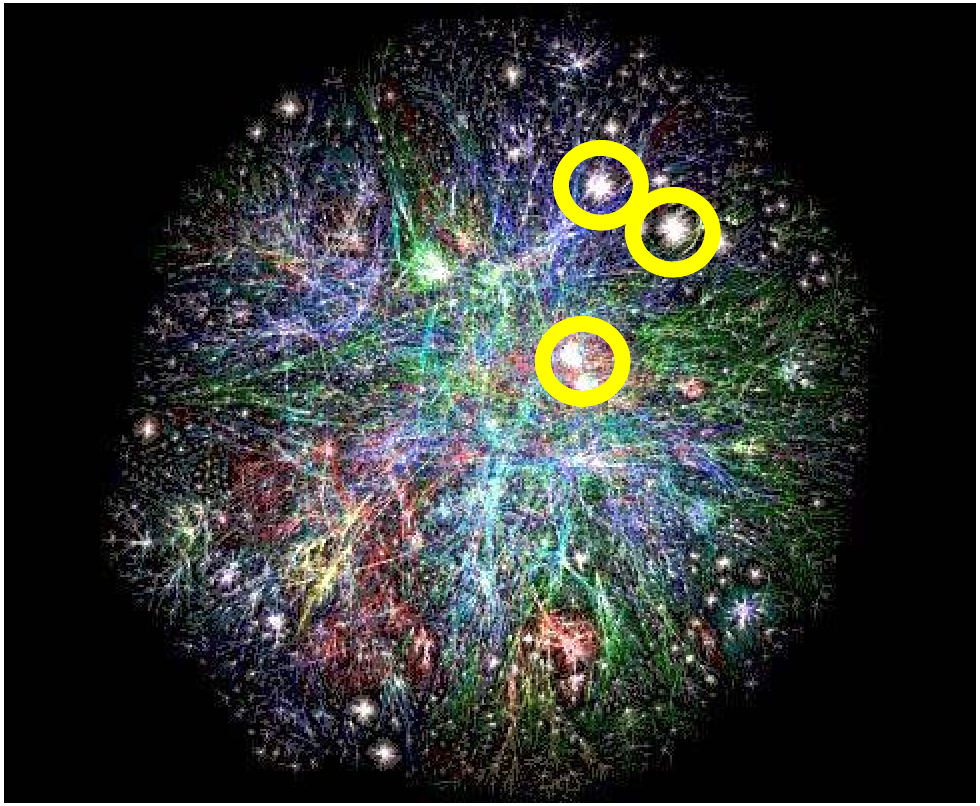

Note the strong local clustering

This is not a random graph

- Graph layout is difficult

[http://img.webme.com/pic/c/chegga-hp/opte_org.jpg]

Universities Linking to Universities

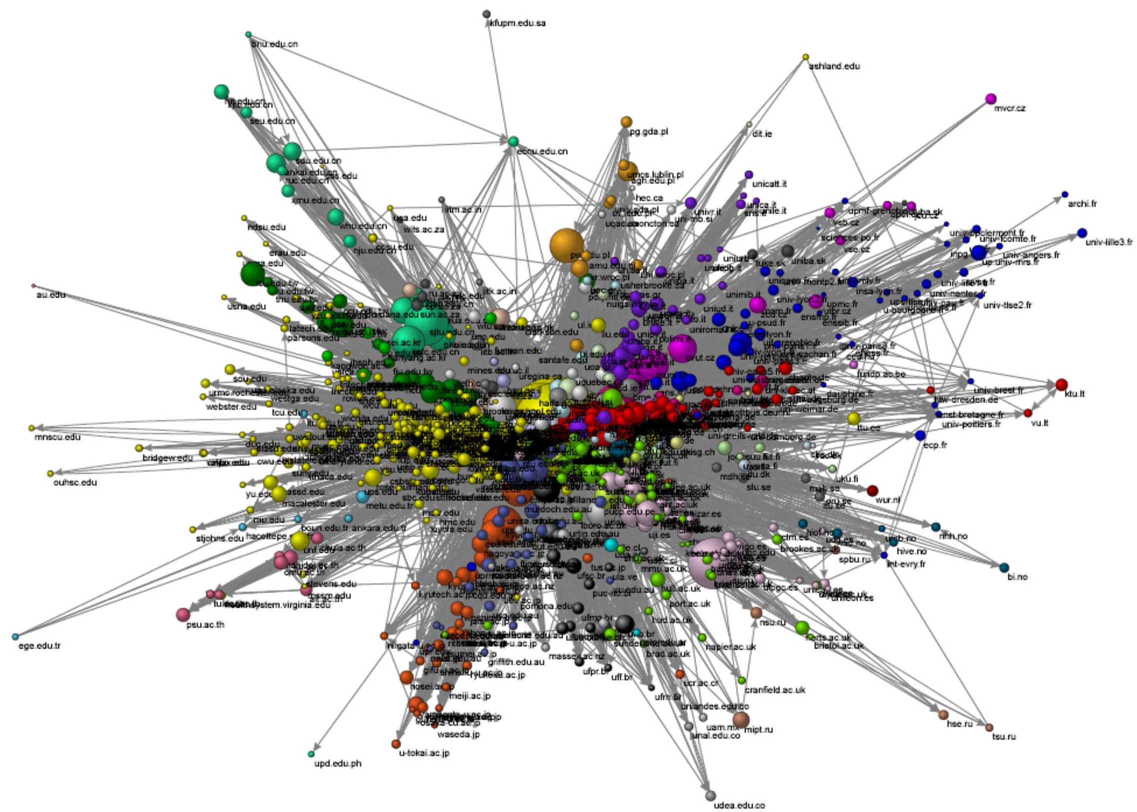

- Small-World Property

[http://internetlab.cindoc.csic.es/cv/11/world_map/map.html]

Human Protein-Protein-Interaction Network

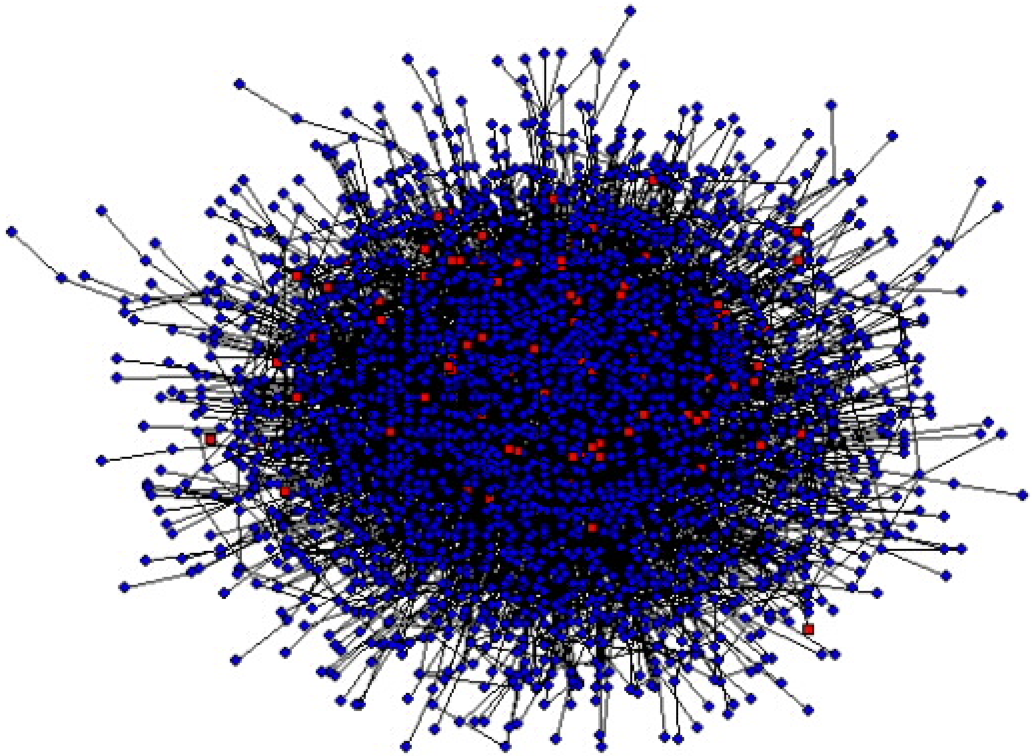

- Still terribly incomplete
- Proteins that are close in the graph likely share function

[http://www.estradalab.org/research/index.html]

Word Co-Occurrence

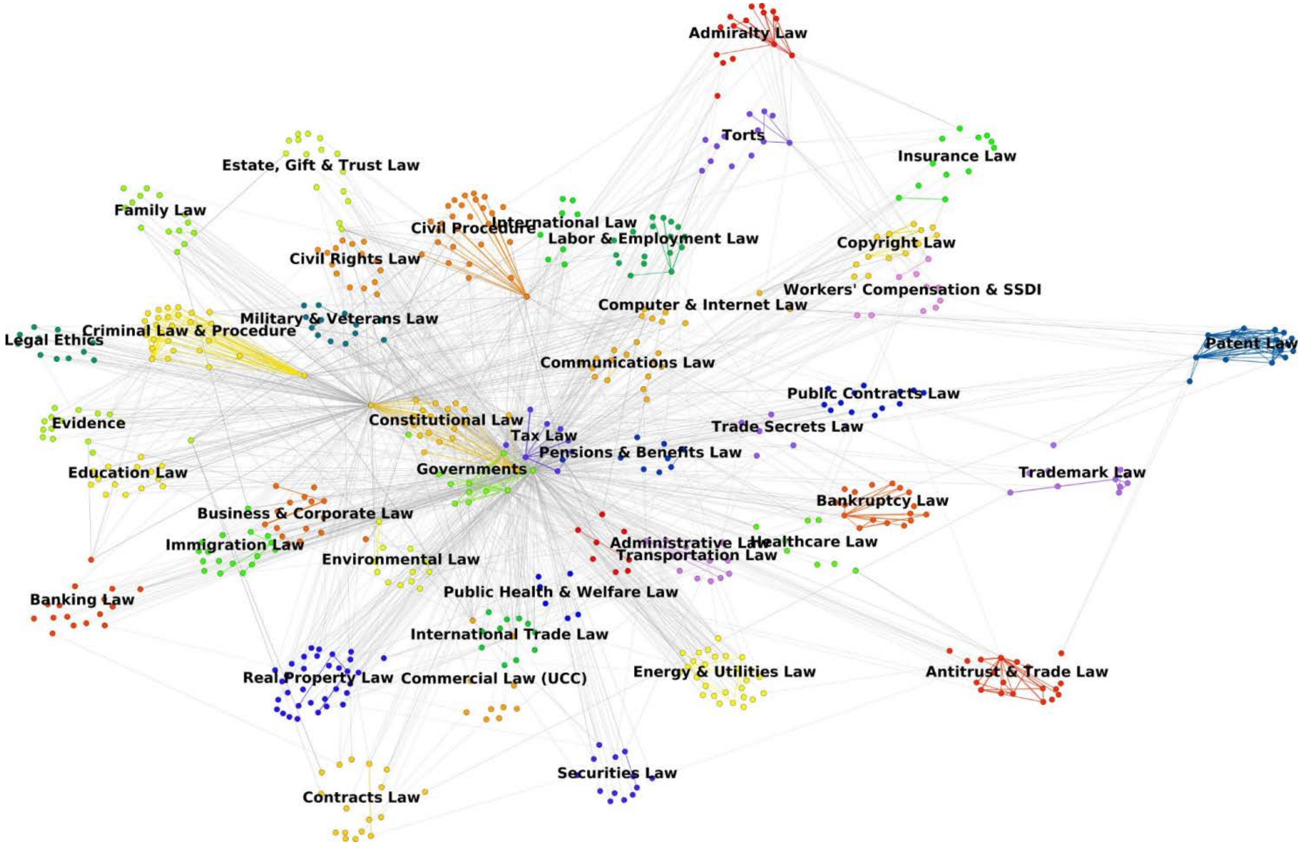

- Words that are close have similar meaning
- Words cluster into topics

[http://www.michaelbommarito.com/blog/]

Social Networks

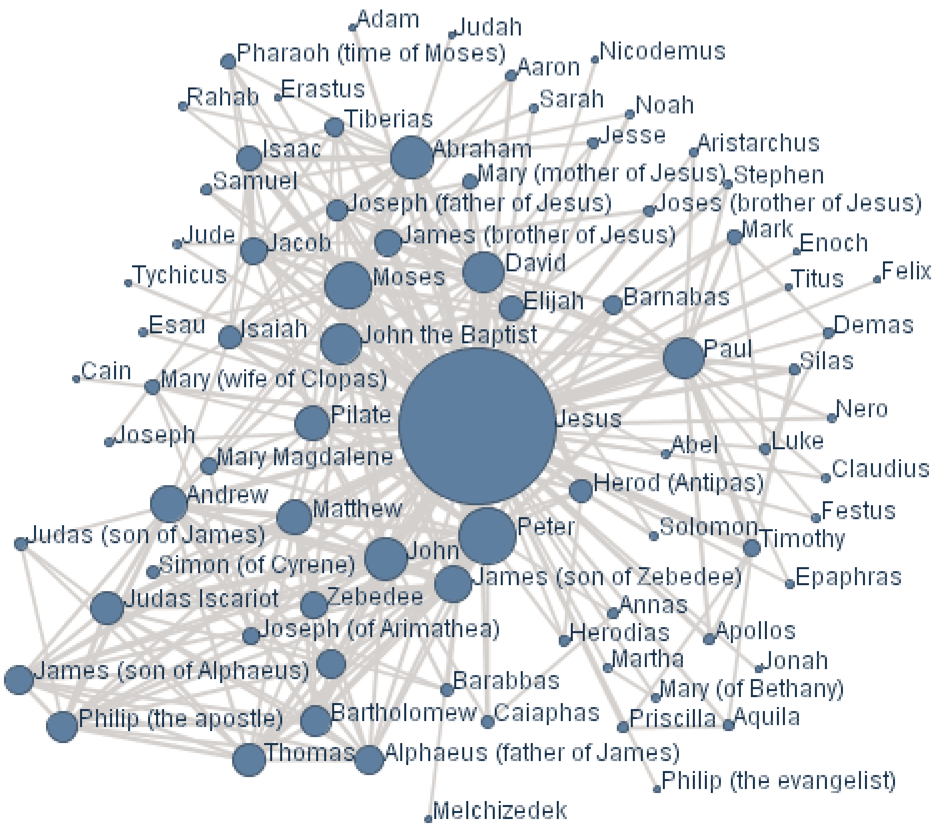

- Six degrees of separation

[http://tugll.tugraz.at/94426/files/-1/2461/2007.01.nt.social.network.png]

Road Network

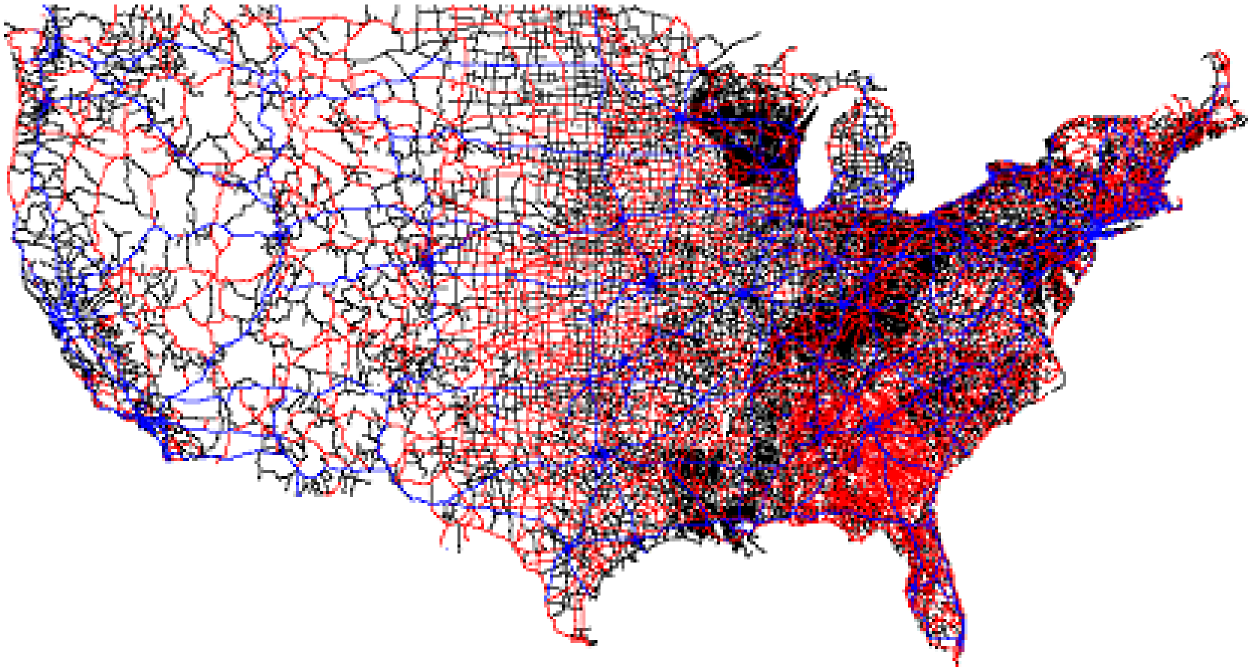

- Specific property: Planar graphs

[Sanders, P. &Schultes, D. (2005).Highway Hierarchies Hasten Exact Shortest Path Queries. In *13th European Symposium on Algorithms (ESA), 568-579.*]

More Examples

- Graphs are also a wonderful abstraction

Coloring Problem

- How many colors do one need to color a map such that never two colors meet at a border?

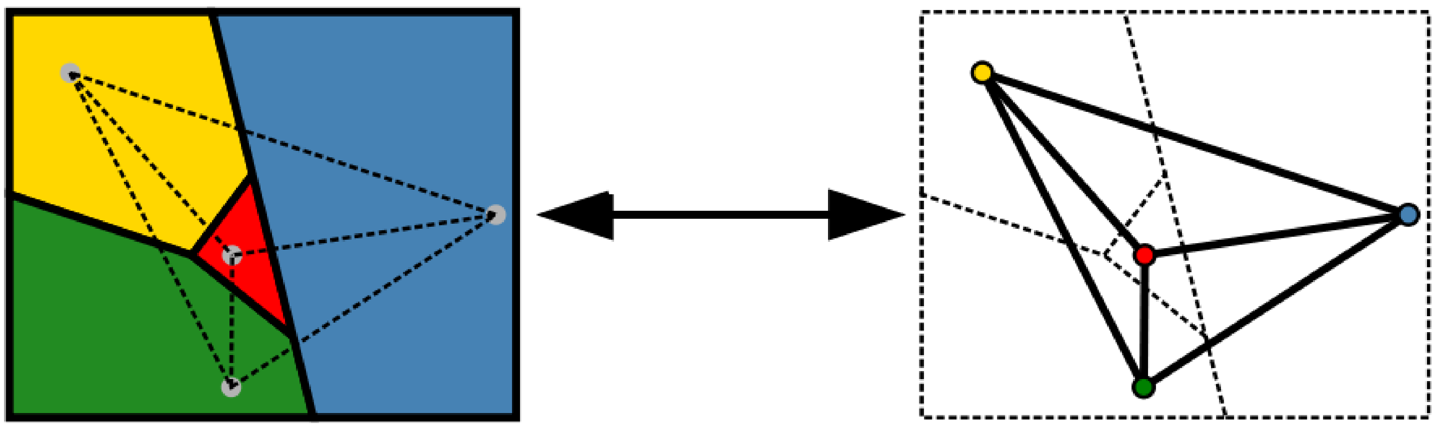

[http://www.wikipedia.de]

- Chromatic number: Number of colors sufficient to color a graph such that no adjacent nodes have the same color
- Every planar graph has chromatic number of at most 4

History [Wikipedia.de]

- This is not simple to proof
- It is easy to see that one sometimes needs at least four colors
- It is easy to show that one may need arbitrary many colors for general graphs
- First conjecture which until today was proven only by computers
  - Falls into many, many subcases – try all of them with a program

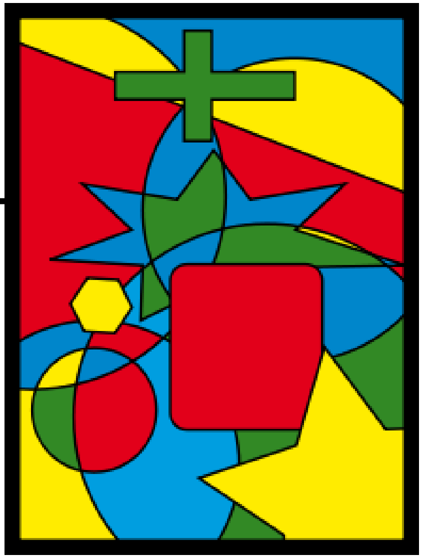

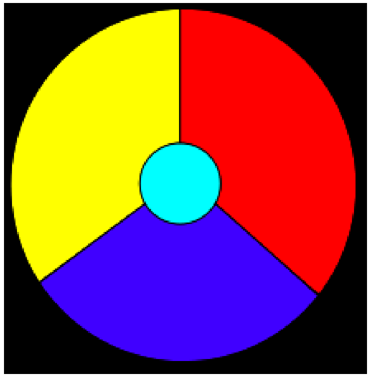

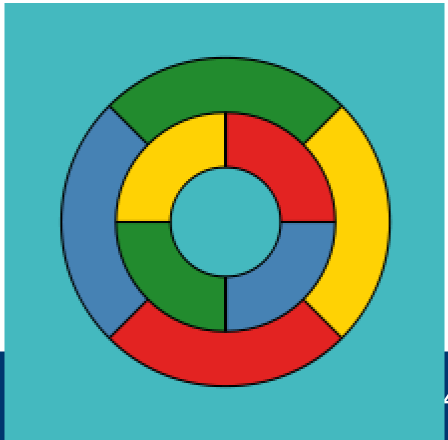

Königsberger Brückenproblem

- Given a city with rivers and bridges: Is there a cycle-free path crossing every bridge exactly once?
  - Euler-Path

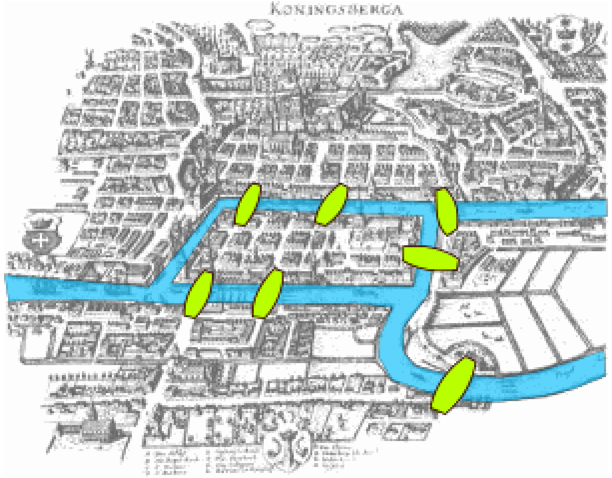

Source: Wikipedia.de

Königsberger Brückenproblem

- Given a city with rivers and bridges: Is there a cycle-free path crossing every bridge exactly once?
  - Euler-Path – simple
- Hamiltonian path
  - ... visits each vertex exactly once
  - NP complete

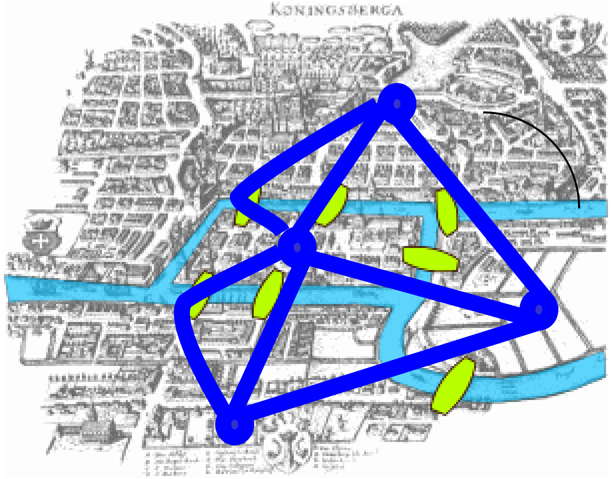

Recall?

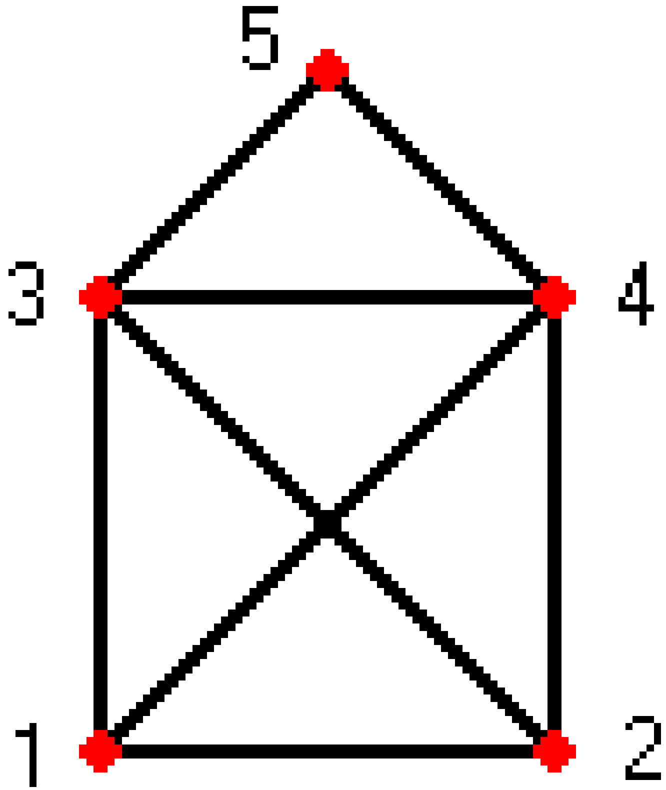

Content of this Lecture

- Graphs
- Representing Graphs
- Traversing Graphs
- Connected Components
- Shortest Paths

Recall from Trees

- Definition
  A graph $G=(V, E)$ consists of a set of vertices (nodes) V and a set of edges ($E \subseteq V \times V$).
  - *A sequence of edges $e_1, e_2, .., e_n$ is called a path iff $\forall 1 \leq i < n$: $e_i=(v', v)$ and $e_{i+1}=(v, v'')$; the length of this path is n*
  - *A path $(v_1,v_2), (v_2,v_3), ..., (v_{n-1},v_n)$ is acyclic iff all v_i are different*
  - *G is acyclic, if no path in G contains a cycle; otherwise it is cyclic*
  - *A graph is connected if every pair of vertices is connected by at least one path*
- Definition
  A graph (tree) is called undirected, if $\forall (v,v') \in E \Rightarrow (v',v) \in E$. Otherwise it is called directed.

More Definitions

- Definition
  Let $G=(V, E)$ be a directed graph. Let $v \in V$
  - *The outdegree out(v) is the number of edges with v as start point*
  - *The indegree in(v) is the number of edges with v as end point*
  - *G is edge-labeled, if there is a function $w: E \rightarrow L$ that assigns an element of a set of labels L to every edge*
  - *A labeled graph with $L=\mathbb{N}$ is called weighted*
- Remarks
  - Weights can as well be reals; often we only allow positive weights
  - Labels / weights max be assigned to edges or nodes (or both)

Some More Definitions

- Definition. *Let $G=(V, E)$ be a directed graph.*
  - *Any $G'=(V', E')$ is called a subgraph of G, if $V' \subseteq V$ and $E' \subseteq E$ and for all $(v_1,v_2) \in E'$: $v_1,v_2 \in V'$*
  - *For any $V' \subseteq V$, the graph $(V', E \cap (V' \times V'))$ is called the induced subgraph of G (induced by V')*

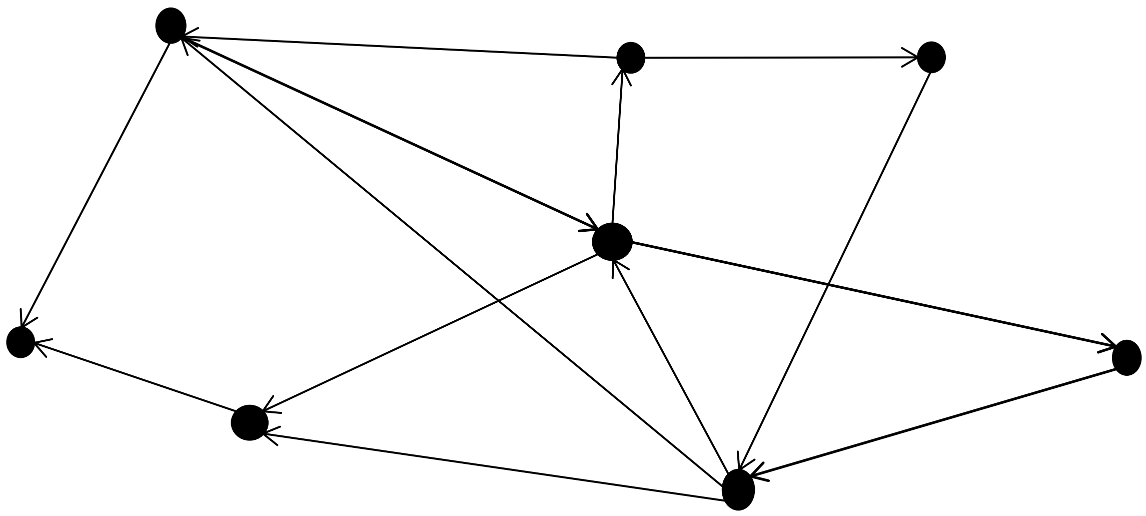

Some More Definitions

- Definition. *Let G=(V, E) be a directed graph.*
  - *Any G'=(V', E') is called a subgraph of G, if $V' \subseteq V$ and $E' \subseteq E$ and for all $(v_1,v_2) \in E'$: $v_1,v_2 \in V'$*
  - *For any $V' \subseteq V$, the graph $(V', E \cap (V' \times V'))$ is called the induced subgraph of G (induced by V')*

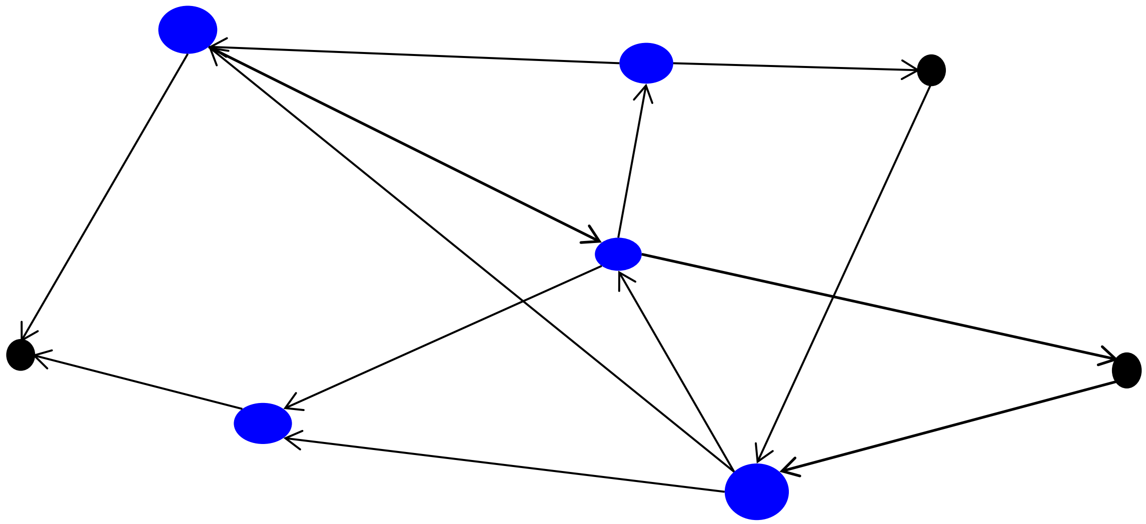

Some More Definitions

- Definition. *Let $G=(V, E)$ be a directed graph.*
  - *Any $G'=(V', E')$ is called a* subgraph of G*, if $V'\subseteq V$ and $E'\subseteq E$ and for all $(v_1,v_2)\in E'$: $v_1,v_2\in V'$*
  - *For any $V'\subseteq V$, the graph $(V', E\cap(V'\times V'))$ is called* the induced subgraph of G *(induced by V')*

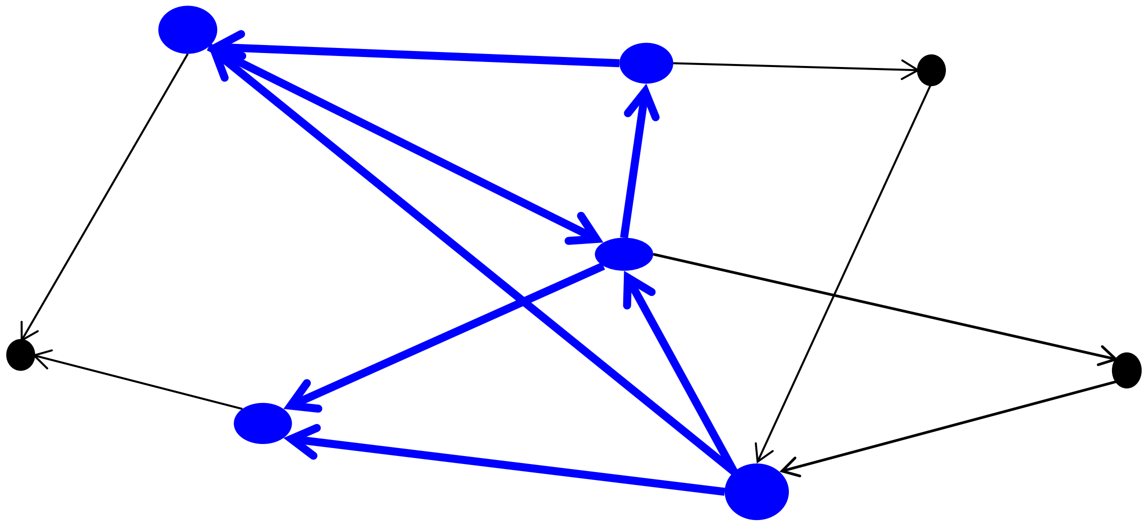

Data Structures

- From an abstract point of view, a graph is a list of nodes and a list of (weighted, directed) edges
- Two fundamental implementations
  - Adjacency matrix
  - Adjacency lists
- As usual, the representation determines which primitive operations take how long
- Appropriateness depends on the specific problem one wants to study and the nature of the graphs
  - Shortest paths, transitive hull, cliques, spanning trees, ...
  - Random, sparse/dense, scale-free, planar, bipartite, ...

Example [OW93]

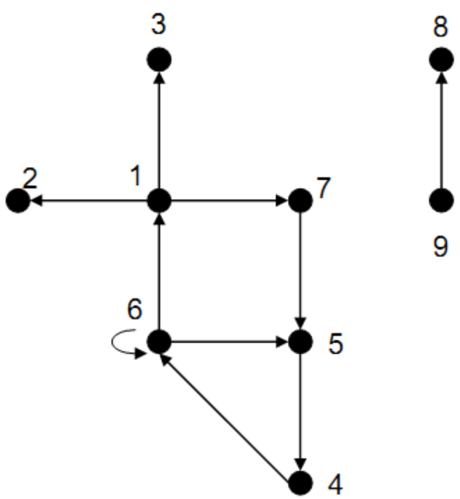

	1	2	3	4	5	6	7	8	9
1	0	1	1	0	0	0	1	0	0
2	0	0	0	0	0	0	0	0	0
3	0	0	0	0	0	0	0	0	0
4	0	0	0	0	0	1	0	0	0
5	0	0	0	1	0	0	0	0	0
6	1	0	0	0	1	1	0	0	0
7	0	0	0	0	1	0	0	0	0
8	0	0	0	0	0	0	0	0	0
9	0	0	0	0	0	0	0	1	0

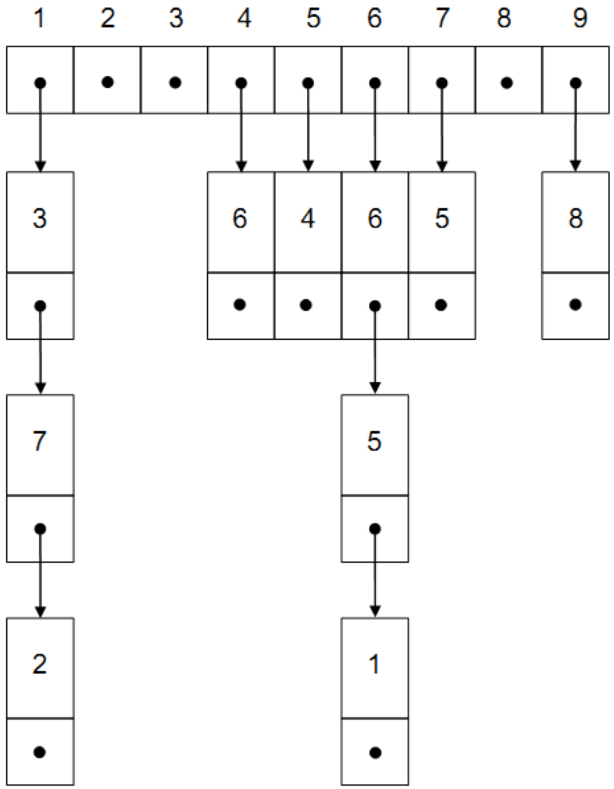

Adjacency Matrix

- Definition
  Let G=(V, E) be a simple graph. The adjacency matrix M_G for G is a two-dimensional matrix of size |V||V|, where $M[i,j]=1$ iff $(v_i,v_j)\in E$*
- Remarks
  - Allows to test existence of an edge in O(1)
  - Requires O(|V|) to obtain all incoming (outgoing) edges of a node
  - For large graphs, M is too large to be of practical use
  - If G is sparse (much less edges than $|V|^2$), M wastes a lot of space
  - If G is dense, M is a very compact representation (1 bit / edge)
  - In weighted graphs, M[i,j] contains the weight
  - Since M must be initialized with zero's, without further tricks all algorithms working on adjacency matrices are in $\Omega(|V|^2)$

Adjacency List

- Definition
  Let $G=(V, E)$. The adjacency list L_G for G is a list containing all nodes of G. The entry representing $v_i \in V$ also contains a list of all edges outgoing (or incoming or both) from v_i.
- Remarks
  – Let k be the maximal outdegree of G. Then, accessing an edge is in O(log(k)) if the edge lists are sorted (or use hashing)
    - Which means O(log(|V|)) in the worst case (for simple graphs)
  – Obtaining a list of all outgoing edges from a node is in O(k)
    - If only outgoing edges are stored, obtaining a list of all incoming edges is O(|V|*log(|E|)) – we need to search all lists
    - Therefore, usually outgoing and incoming edges are stored, which considerably increases space consumption
  – If G is sparse, L is a compact representation

Comparison

	M	**L**
Test an edge	O(1)	O(log(k))
All outgoing edges of a node	O(n)	O(k)
Space	$O(n^2)$	O(n+m)

- With n=|V|, m=|E|
- We assume a node-indexed array / a node-index list
  - L is an array and nodes are unique numbered

Transitive Closure

- Definition
  Let $G=(V,E)$ be a digraph and $v_i,v_j \in V$. The *transitive closure* *of G is a graph $G'=(V, E')$ where $(v_i,v_j) \in E'$ iff G contains a path from v_i to v_j.*
- TC usually is represented as adjacency matrix

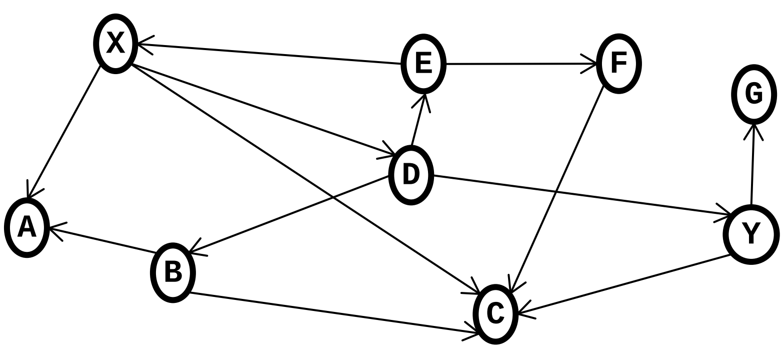

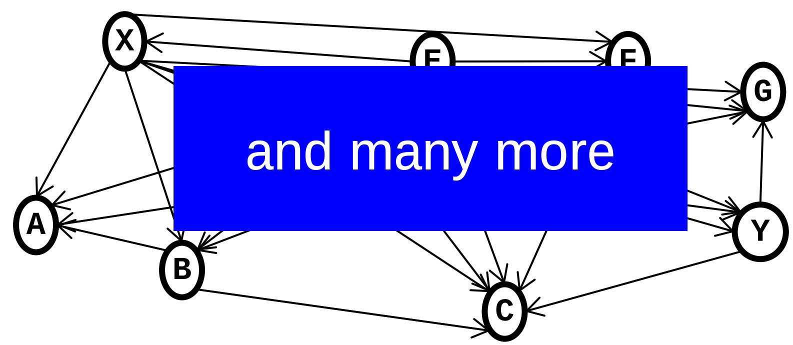

Content of this Lecture

- Graphs
- Representing Graphs
- Traversing Graphs
- Connected Components
- Shortest Paths

Graph Traversal

- One thing we often do with graphs is traversing them
- “Traversal” means visiting every node exactly once
  - Not necessarily on one consecutive path (Hamiltonian path)
- Two popular orders
  - Depth-first: Using a stack
  - Breadth-first: Using a queue
  - The scheme is identical to that in tree traversal
- Difference
  - We have to take care of cycles
  - No root – where should we start?

Breaking Cycles

- In a cyclic graph, naïve traversal will ...
  - run into infinite loops: Algorithm does not terminate
  - visit nodes more than once
- Breaking cycles / avoiding multiple visits
  - Assume we started the traversal at a node r
  - During traversal, we kept a list S of already visited nodes and are now in node v and aim to proceed to v' using e
    - Because e=(v, v')∈E and e was not used before from v
  - If v'∈S, v' was visited before and we are about to run into a cycle
  - In this case, e is ignored

Example

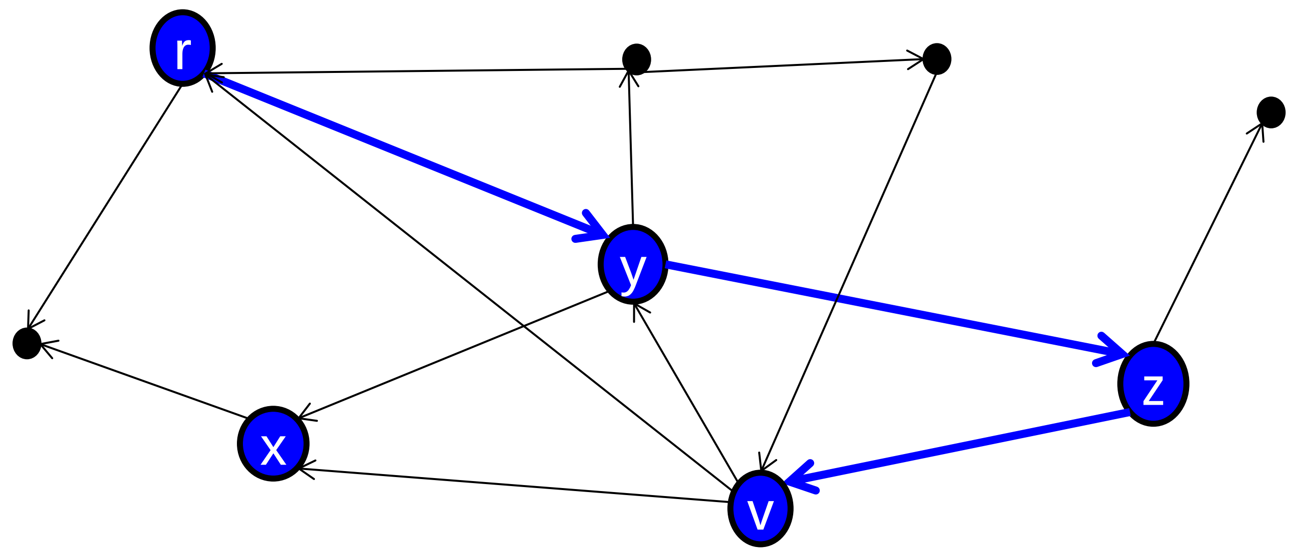

- Started at r and went S={r, y, z, v}
- Testing (v,y): y∈S, drop
- Testing (v, r): r∈S, drop
- Testing (v, x): x∉S, proceed

Where do we Start?

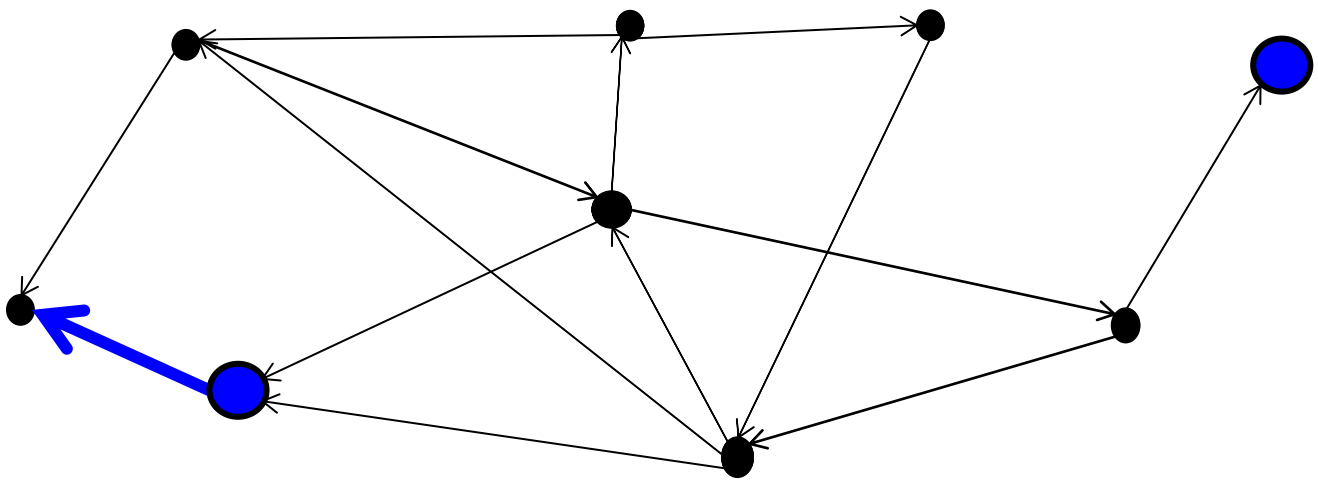

Where do we Start?

- Definition
  Let G=(V, E). Let V'⊆V and G' be the subgraph of G induced by V'
  - *G' is called connected if it contains a path between any pair v,v'∈V'*
  - *G' is called maximally connected, if no subgraph induced by a superset of V'* is connected*
  - *Any maximal connected subgraph of G is called a connected component of G, if G is undirected, and a strongly connected component, if G is directed*

Example

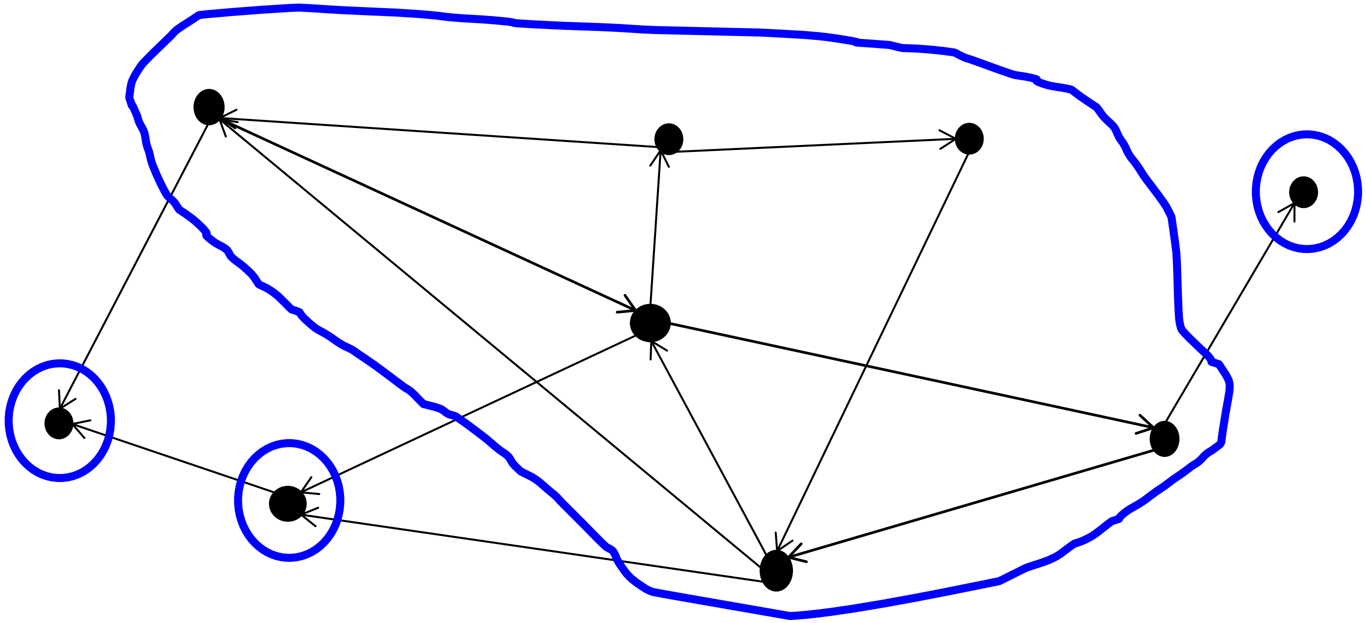

Where do we Start?

- If a undirected graph falls into several connected components, we cannot reach all nodes by a single traversal, no matter which node we use as start point
- If a digraph falls into several strongly connected components, we might not reach all nodes by a single traversal
- Remedy: We restart at unseen nodes until all nodes have been traversed

Depth-First Traversal on Graphs

```
func void DFS ((V,E) graph) {
  U := V;        # Unseen nodes
  S := Ø;        # Seen nodes
  while U≠∅ do
    v := any_node_from( U);
    traverse( v, S, U);
  end while;
}
```

```
func void traverse (v node,
                    S,U list)
{
  s := new Stack();
  s.put( v);
  while not s.isEmpty() do
    n := s.get();
    print n;   # Do something
    U := U \ {n};
    S := S ∪ {n};
    c := n.outgoingNodes();
    foreach x in c do
      if x∈U then
        s.put( x);
      end if;
    end for;
  end while;
}
```

Called once for every connected component

Analysis

- We have every node exactly once on the stack
  - Once visited, never visited again
- We look at every edge exactly once
  - Outgoing edges of every visited node are never considered again
- S and U can be implemented as bit-array of size |V|, allowing O(1) operations
- Altogether: O(n+m)

```
func void traverse (v node,
                    S,U list) {
  s := new Stack();
  s.put( v);
  while not s.isEmpty() do
    n := s.get();
    print n;
    U := U \ {n};
    S := S ∪ {n};
    c := n.outgoingNodes();
    foreach x in c do
      if x∈U then
        s.put( x);
      end if;
    end for;
  end while;
}
```

Unusual Traversal: Random Surfer

- How do search engines determine which hits appear first?
- Ingredient 1: Match to the query
- Ingredient 2: Popularity of the page
  – Pages with many incoming edges are popular
  – Pages are the more popular, the more popular the start points of incoming edges are
  – Recursive definition ...
- Random surfer model
  – Assume a surfer starting at a page chosen at random and following links at random for an infinite time
    - May jump to random pages if stuck
  – Which fraction of time will he spend in a given page?

Content of this Lecture

- Graphs
- Representing Graphs
- Traversing Graphs
- Connected Components
- Shortest Paths

In Undirected Graphs

- In an undirected graph, whenever there is a path from r to v and from v to v', then there is also a path from v' to r
  - Simply go the path r $\rightarrow$ v $\rightarrow$ v' backwards
- Thus, DFS (and BFS) traversal can be used to find all connected components of a undirected graph G
  - Whenever you call traverse(v), create a new component
  - All nodes visited during traverse(v) are added to this component
- Obviously in O(n+m)

In Digraphs

- The problem is considerably more complicated for digraphs
  - Previous conjecture does not hold
- Still: Tarjan‘s or Kosaraju’s algorithm find all strongly connected components in O(n + m)
  - See next lecture

Content of this Lecture

- Graphs
- Representing Graphs
- Traversing Graphs
- Connected Components
- Shortest Paths
  - Single-Source-Shortest-Paths: Dijkstra’s Algorithm
  - Shortest Path between two nodes
  - Other

Distance in Graphs

- Definition
  Let G=(V, E) be a graph. The distance d(u,v) between any two nodes u and v from V is defined as
  - *G un-weighted: The length of the shortest path from u to v, or ∞ if no path from u to v exists*
  - *G weighted: The minimal aggregated edge weight of all non-cyclic paths from u to v, or ∞ if no path from u to v exists*
- Remark
  - Distance in un-weighted graphs is the same as distance in weighted graphs with unit costs
  - Beware of negative cycles in directed graphs

Single-Source Shortest Paths in a Graph

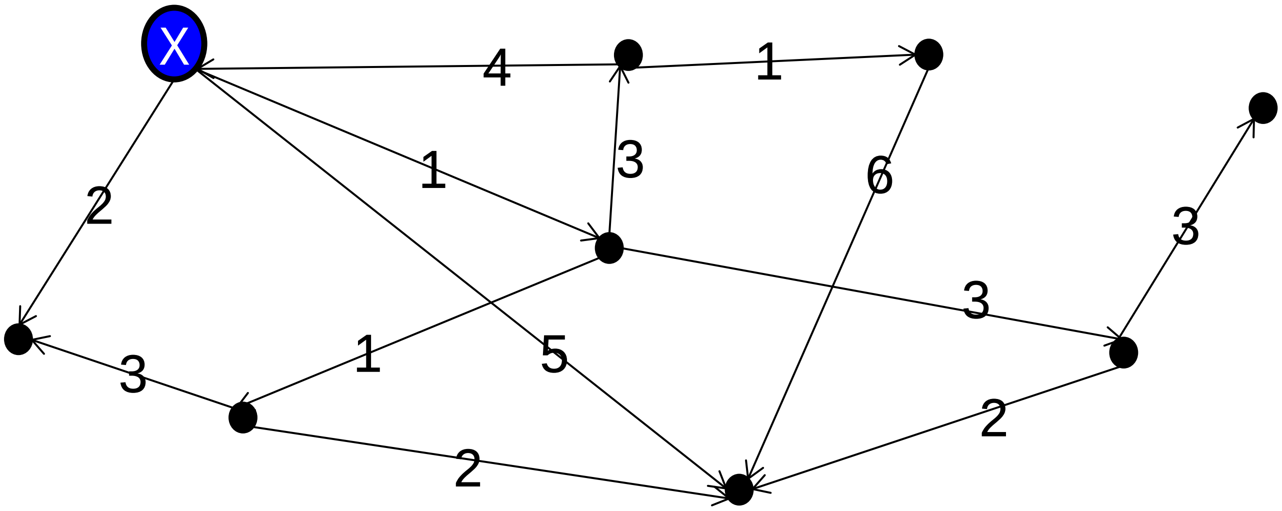

- Task: Find the distance between X and all other nodes
  - Here: Only positive edge weights (see next lecture)

Dijkstra's algorithm

```
1.  G = (V, E);
2.  x : start_node;    # x∈V
3.  A : array_of_distances_from_x;
4.  ∀i: A[i]:= ∞;
5.  L := V;       # organized as PQ
6.  A[x] := 0;
7.  while L≠∅
8.    k := L.get_closest_node();
9.    L := L \ k;
10.   forall (k,f,w)∈E do
11.     if f∈L then
12.       new_dist := A[k]+w;
13.       if new_dist < A[f] then
14.         A[f] := new_dist;
15.       end if;
16.       update( L);
17.     end if;
18.   end for;
19. end while;
```

- Assume a heap-based PQ L
- L holds at most all nodes (n)
- L4: O(n)
- L5: O(n*log(n)) (build PQ)
- L8: O(1) (getMin)
- L9: O(log(n)) (deleteMin)
- L10: O(m) (with adjacency list)
- L11: O(1)
  - Requires additional array of nodes
- L16: O(log(n)) (updatePQ)

Dijkstra’s algorithm

```
1.  G = (V, E);
2.  x : start_node;    # x∈V
3.  A : array_of_distances;
4.  ∀i: A[i]:= ∞;
5.  L := V;      # organized as PQ
6.  A[x] := 0;
7.  while L≠∅
8.    k := L.get_closest_node();
9.    L := L \ k;
10.   forall (k,f,w)∈E do
11.     if f∈L then
12.       new_dist := A[k]+w;
13.       if new_dist < A[f] then
14.         A[f] := new_dist;
15.       end if;
16.       update( L);
17.     end if;
18.   end for;
19. end while;
```

- Central costs
  - L9: O(log(n)) (deleteMin)
  - L10: O(m) (adjacency list)
  - L16: O(log(n))
- Loops
  - Lines 7-18: O(n)
  - Line 10-17: All edges exactly once
  - Together: O(m+n)
- Altogether: O((m+n)*log(n))
  - Also possible in $O(n^2)$; better in very dense graphs ($m \sim n^2$)

Single-Source, Single-Target

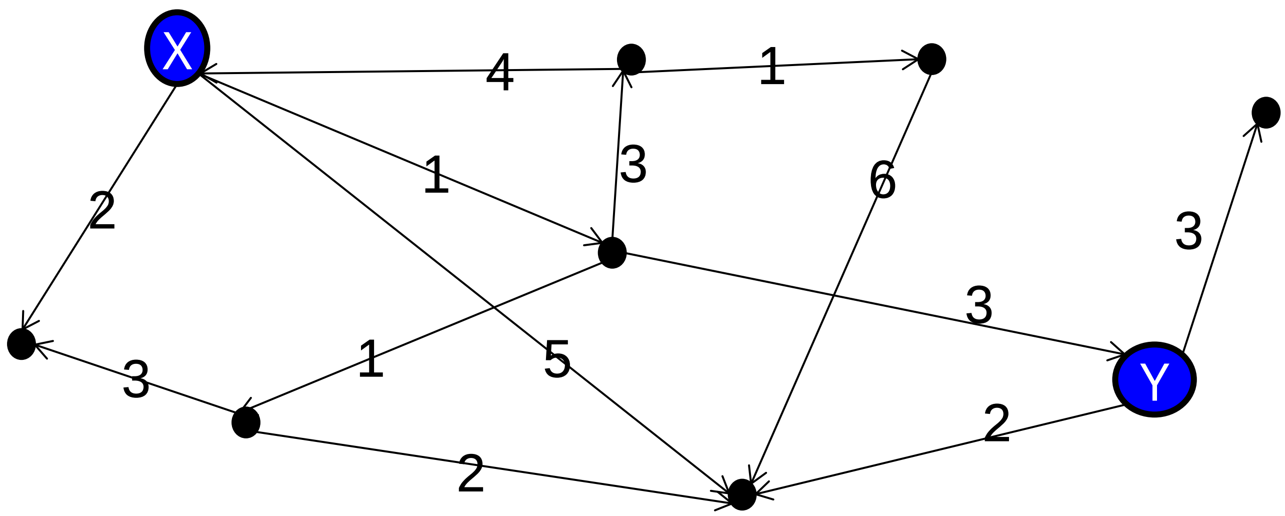

- Task: Find the distance between X and only Y
  - In general, there is no way to be WC-faster than Dijkstra
    - We can stop as soon as Y appears at the min position of the PQ
    - We can visit edges in order of increasing weight
    - Worst-case complexity unchanged, average case is (slightly) better
- Things are different in planar graphs (navigators!)

Faster SS-ST Algorithms

- Trick 1: Pre-compute all distances
  - Transitive closure with distances
  - Requires $O(|V|^2)$ space: Prohibitive for large graphs
  - How? See next lecture

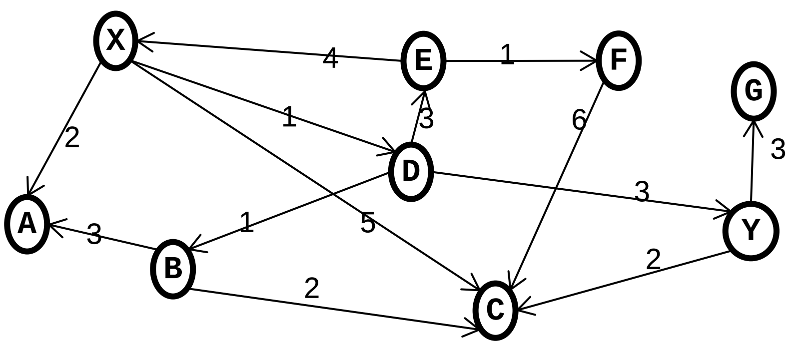

→	A	B	C	D	E	F	G	X	Y
A	0	-	-	-	-	-	-	-	-
B	3	0	2	-	-	-	-	-	-
C	-	-	0	-	-	-	-	-	-
D	4	1	3	0	3	4	6	7	3
E	6	6	7	5	0	1	11	4	8
F	-	-	6	-	-	0	-	-	-
G	-	-	-	-	-	-	0	-	-
X	2	2	4	1	4	5	7	0	4
Y	-	-	2	-	-	-	3	-	0

Faster SS-ST Algorithms

- Trick 2: Two-hop cover with distances
  - Find a small set S of nodes such that
    - For every pair of nodes v_1,v_2, at least one shortest path from v_1 to v_2 goes through a node $s \in S$
    - Thus, the distance between v_1,v_2 is $\min\{ d(v_1,s)+d(s,v_2) \mid s \in S)$
    - S is called a 2-hop cover
  - Problem: Finding a minimal S is NP-complete
    - And S need not be small

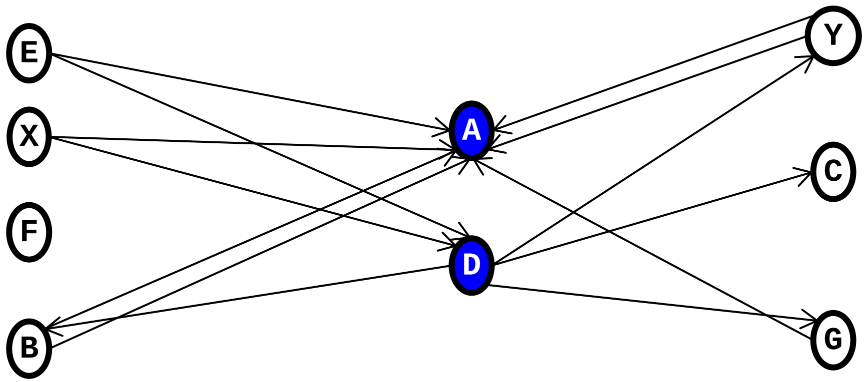

More Distances

- Graphs with negative edge weights
  - Shortest paths (in terms of weights) may be very long (in terms of edges)
  - Bellman-Ford algorithm in $O(n^2*m)$
- All-pairs shortest paths
  - Only positive edge weights: Use Dijkstra n times
  - With negative edge weights: Floyd-Warshall in $O(n^3)$
    - See next lecture
- Reachability
  - Simple in undirected graphs: Compute all connected components
  - In digraphs: Use Dijkstra or a special graph indexing method
    - See special modules